


Introduction to Serial Code Optimisations

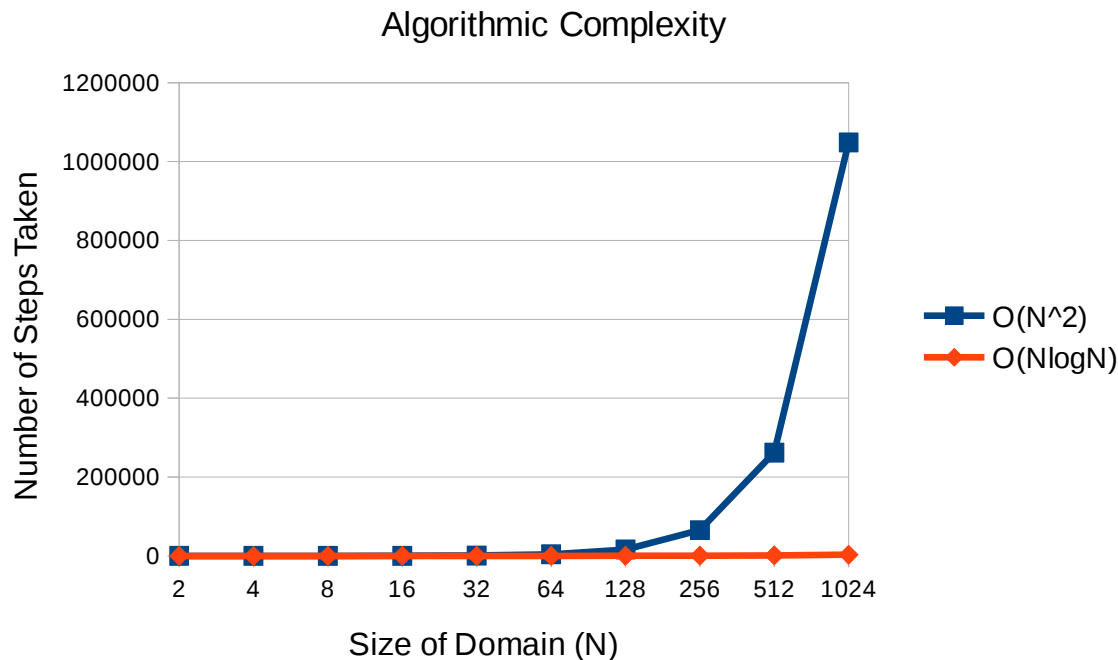


Getting from A to B – faster!

Overview

- Some hints and tips:
 - Lesson #1: Algorithms matter.
 - Lesson #2: Focus your efforts.
 - Never guess, use a profiler instead.
 - Lesson #3: Experiment with compiler flags ***and then re-profile.***
 - Lesson #4: The memory hierarchy has a large effect (and it grows over time).
 - Lesson #5: Making use of wide registers is important on modern processors.

Lesson #1: Algorithms Matter



A solution in $O(N^2)$..

What do you mean?


Let's say we have a problem 'of size N'. This could be sorting a list of N items.

Sorting algorithm A takes $O(N^2)$ operations to sort the list. Algorithm B takes $O(N \log N)$ operations..

- Algorithm A $\rightarrow O(N^2)$.
- Algorithm B $\rightarrow O(N \log N)$.
- For large N, even a lousy implementation of B will beat a great implementation of A.

Lesson #2: Focus your Efforts

"We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil**. Yet we should not pass up our opportunities in that critical 3%. A good programmer will not be lulled into complacency by such reasoning, he will be wise to look carefully at the critical code; but **only after that code has been identified**" — Donald Knuth



For the assignments, we're going to
(a) assume that we have the best algorithm; and
(b) follow Knuth's sage advice:

- 1 - profile
- 2 - find the critical code
- 3 - only attempt to optimise the critical code
- 4 - **repeat this cycle!**

Lesson #2: Never Guess

Consider the following 2 loops:

```
/* loop 1 */  
for (i=0; i<N; i++)  
    a[i] = pow(6, 0.35);
```

This is repeated
(and hence wasted)
computation, right?

```
/* loop 2 */  
x = pow(6, 0.35);  
for (i=0; i<N; i++)  
    b[i] = x;
```

gcc -O3

Loop containing invariants

Elapsed time: 0.752673 (s)

Manually 'hoisted' loop

Elapsed time: 0.938888 (s)

Lesson #2: Use a Profiler Instead

This is some output from *gprof*:

Each sample counts as 0.01 seconds.

% cumulative time	self seconds	self seconds	self calls	total us/call	total us/call	name
72.69	25.01	25.01	300000	83.37	83.37	collision
15.03	30.18	5.17	300000	17.23	17.23	propagate
9.13	33.32	3.14	300000	10.47	10.47	total_density

This function accounts for the vast majority of the run time.
So let's focus our (initial) efforts there..

Changing the Compiler Flags (Conditionals in a Loop)

```
const size_t arraySize = 32768;  
std::vector<int> data(arraySize);
```

Fill array with
random numbers

```
// random sequence in range [0,256]  
for (unsigned c = 0; c < arraySize; ++c)  
    data[c] = std::rand() % 256;
```

```
// branches predictable if sorted, otherwise random  
std::sort(data.begin(), data.end());
```

Sort, or not..

```
long long sum = 0;  
for (unsigned i = 0; i < 100000; ++i)  
    { for (unsigned c = 0; c < arraySize; ++c)  
        { if (data[c] >= 128)  
            sum += data[c];  
        }  
    }
```


Conditional triggers
Predictably or randomly

Branch Prediction: Processors and Compiler Flags

				Time (s)
gcc 4.4.3	Core 2	-O2	sorted	3.2
			random	13.7
		-O3	sorted	6.4
			random	15.6
gcc 4.4.6	Xeon E5	-O2	sorted	2.0
			random	10.4
		-O3	sorted	2.6
			random	2.6

- Branching in loops can stall the (deep) pipelines in modern processors. (Yet another attempt to find parallelism in code execution)
- Branch prediction clearly matters.
- **BUT..** the details vary from processor to processor and compiler to compiler.

Lesson #4: The Memory Hierarchy has a Large Effect

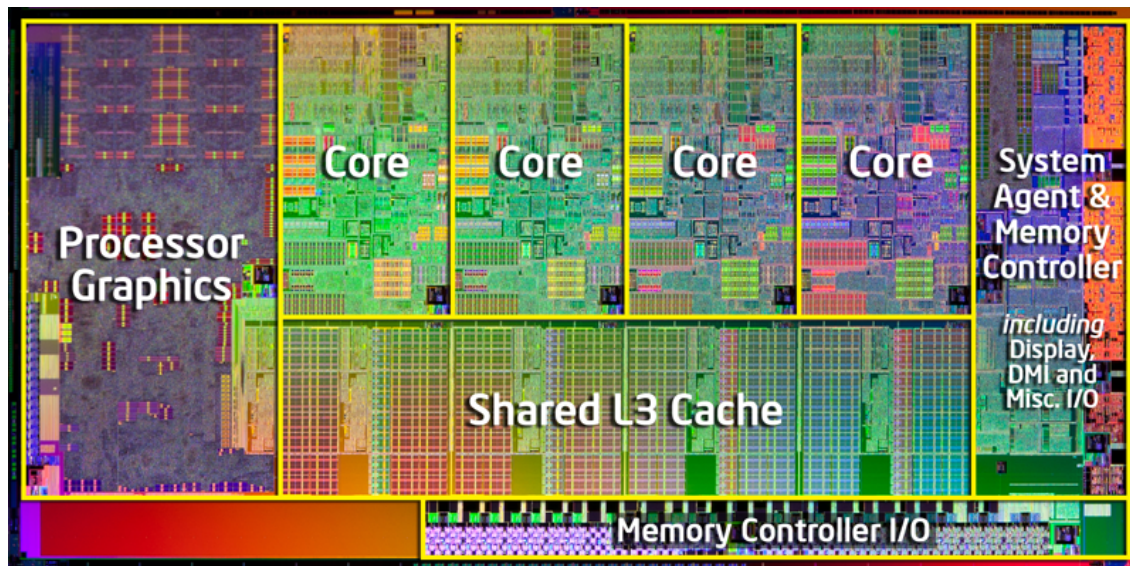



See e.g. the STREAM benchmark:
<https://www.cs.virginia.edu/stream/>

Caches try to Bridge the Gap

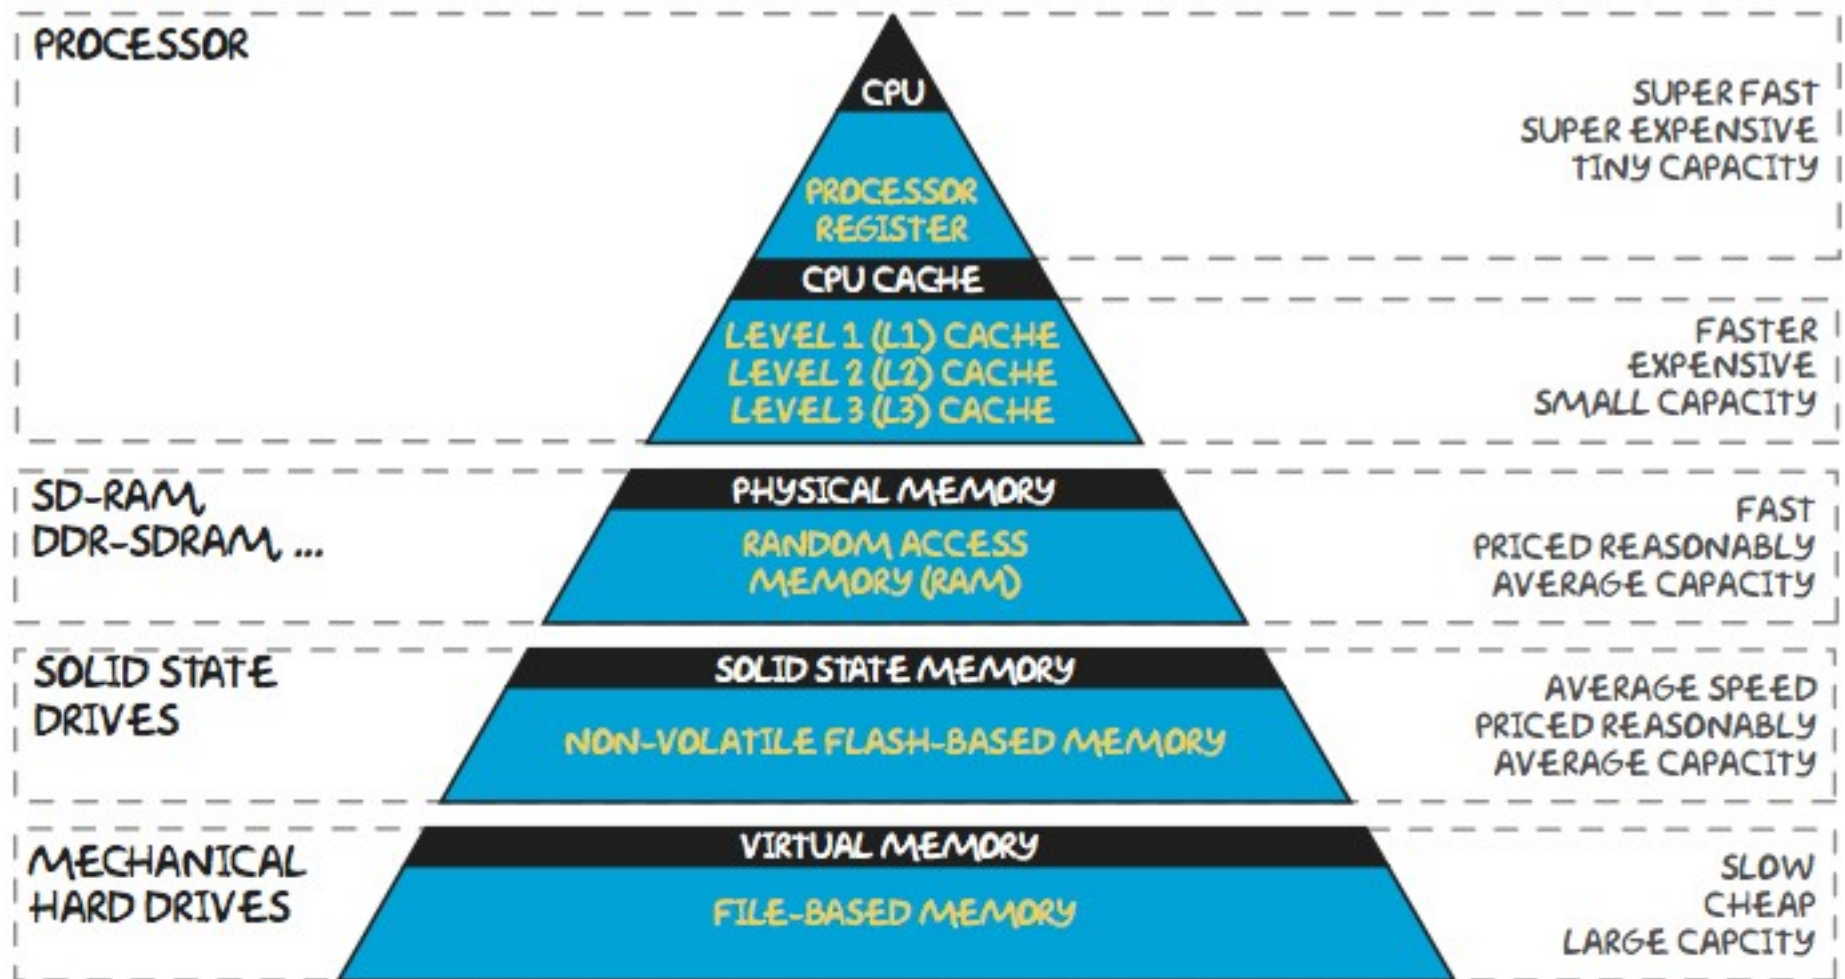
General Layout:

- Registers and Caches on chip
- Bus connection to RAM
- Hard drives slower access times



Schematic for
Intel SandyBridge

THE MEMORY HIERARCHY



Memory Hierarchies: Analogy

- L1 cache: Like picking up a paper from your desk (~3s)
- L2 cache: Like getting up and going to the bookshelf (~15s)
- Main memory: Like walking down the corridor (several minutes)
- Disk?: Like walking around the coastline of Britain (~ 1yr)!




Make the Most of the Cache that You Have

- Explore the use of different data types.
 - Perhaps you can store the same *information* in cache (to a suitable degree of precision), but using less space?
- Make the most of the data that you have in cache already – don't revisit the same data later in your program if you don't have to.

Cache Thrashing: A Synthetic Example


2-way associative
cache



Thrashing is the opposite of re-use:

```
#define max 1024*1024
double a[max], b[max], c[max], d[max]
for(ii=0; ii<max; ii++) {
    a[ii] = b[ii] + c[ii]*d[ii];
}
```

ii



d
c
b

~200 clock
cycles to
load

Accessing **b** will cause
eviction of **c** or **d**

Best to step
contiguously
through cache line

(a block of)
main memory

Row vs. Column Major Order

2D array

A	B	C
D	E	F
G	H	I

Row Major Order

e.g. C

A	B	C	D	E	F	G	H	I
---	---	---	---	---	---	---	---	---

Column Major Order

e.g. Fortran


A	D	G	B	E	H	C	F	I
---	---	---	---	---	---	---	---	---

Row vs. Column Major Order

2D array


A	B	C
D	E	F
G	H	I

Contiguous
Access pattern



This access
pattern is more
efficient

Non-contiguous
Access pattern




Don't Re-Visit Memory More Often Than You Need To

A trivial example:

```
for(ii=0; ii<N; ii++) {  
    grid[ii] = 0;  
}  
  
for(ii=0; ii<N; ii++) {  
    grid[ii] += ...  
}
```

If N is large, we reload the data into cache..



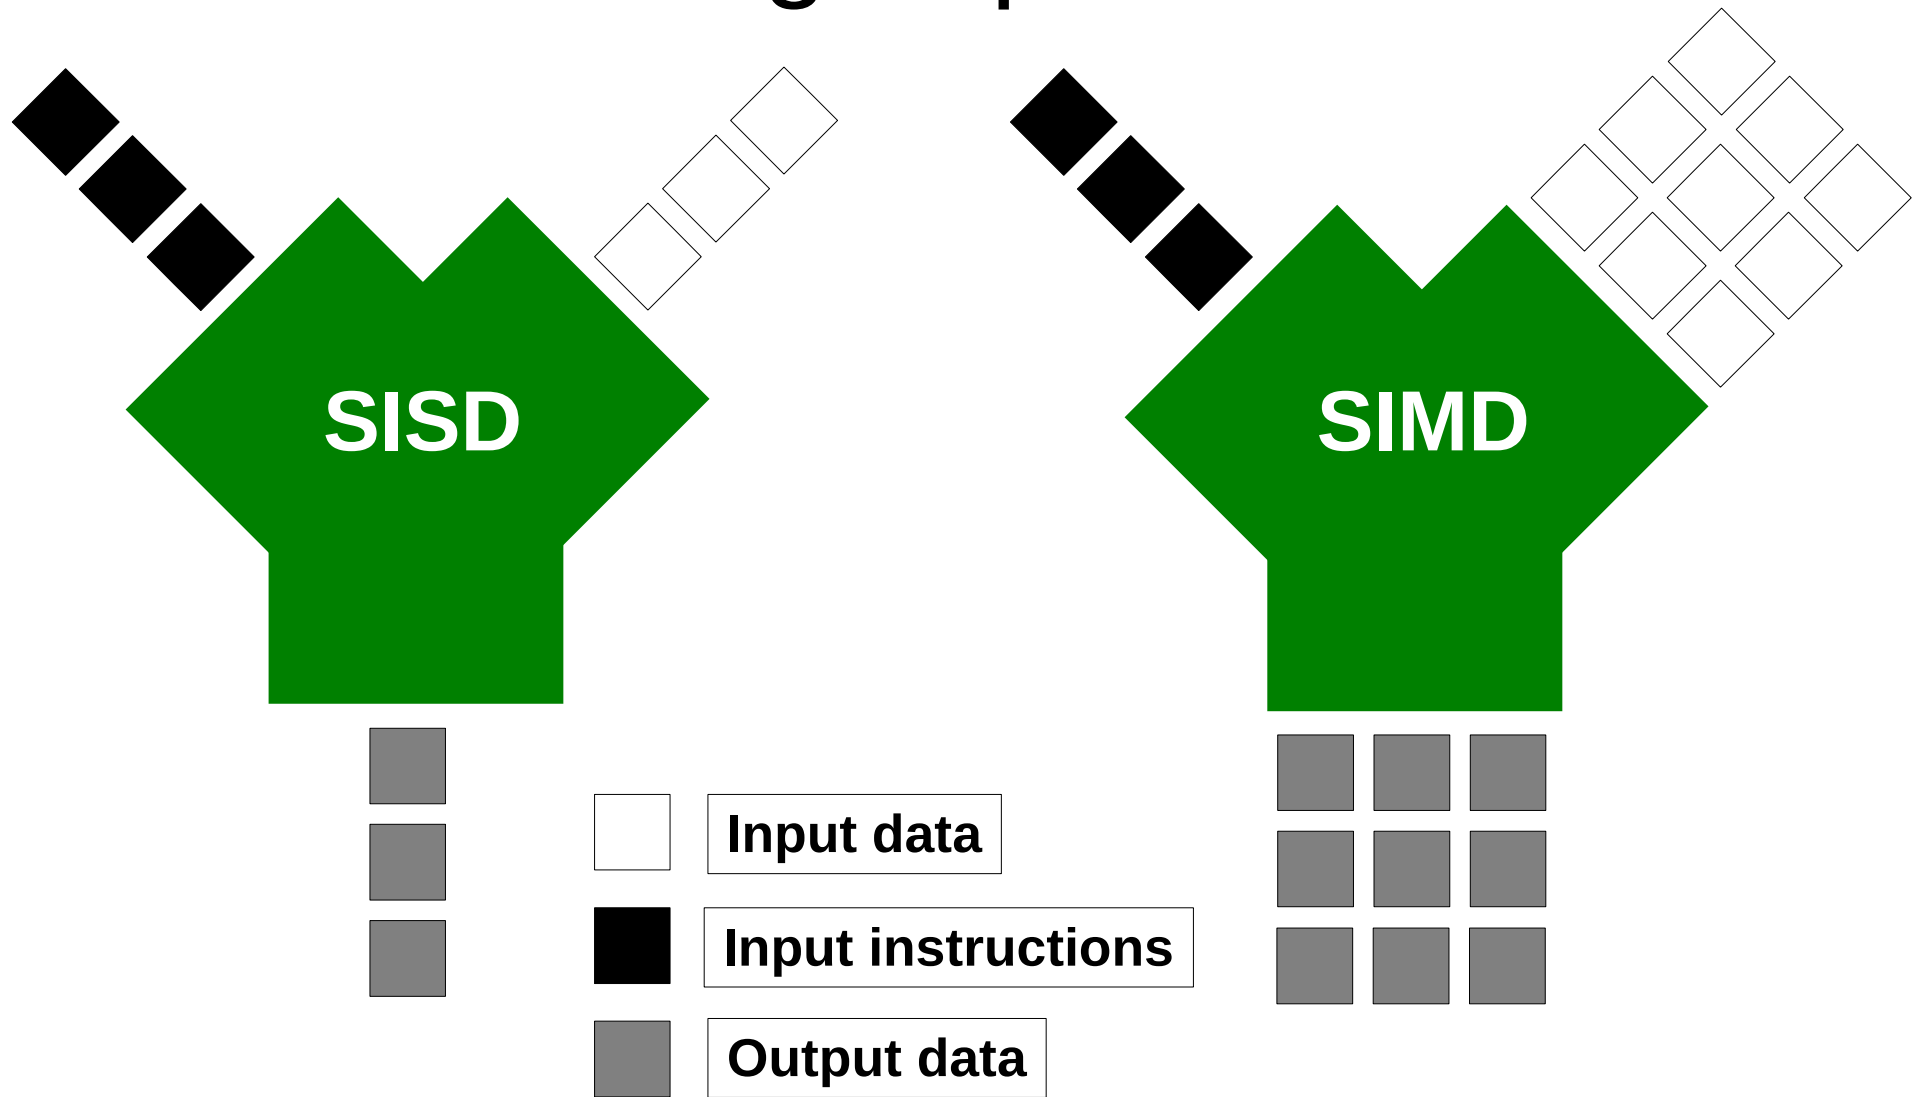
```
for(ii=0; ii<N; ii++) {  
    grid[ii] = 0;  
    grid[ii] += ...  
}
```

merging
loops

Look for similar opportunities
in your assignment.

You may need to re-write the code so that the same operations are performed on the data but with potentially radically different source code.

Lesson #5: Vectorisation is of Growing Importance



Trends in Register Width

- SSE: 128bit, e.g. 2 doubles or 4 floats
 - BCp1 (AMD Opteron 2218)
 - BCp2 (Intel Xeon E5462, 'Harpertown')
- AVX: 256bit, e.g. 4 doubles
 - BCp3 (Intel SandyBridge)
- Intel MIC: 512bit, e.g. 8 doubles
 - BCp3
- Employ through:
 - compiler options - portability.
 - language intrinsics – most control (& hence performance?).

Vectorisation – the Principles

- Memory allocations and data structures aligned to the appropriate boundaries (RAM → caches → registers).
- No data dependencies within loop iterations (e.g. $a[i] = a[i-1] + 1$), as we want to process several iterations at once.
- **Compilers prefer correct programs to fast but incorrect ones!** Watch out for aliased pointers (you may need to use compiler hints, e.g. `#pragma ivdep`, `'restrict'`).
- What function calls are you allowed? Often only inlined and math intrinsics, no switch statements.
- Newer compilers and chipsets offer progressively more support to the programmer: pay attention to the compilers vectorisation reports.

Vectorisation: Many Useful Tutorials

- <https://software.intel.com/en-us/articles/vectorization-essential>
- http://hpac.rwth-aachen.de/teaching/sem-accg-16/slides/08.Schmitz-GGC_Autovec.pdf
- and many others besides.

Appendix – Using gprof

<https://sourceware.org/binutils/docs/gprof/>

'Instrument' the code for profiling..

```
cc -O2 -pg myprog.c -o myprog.exe  
./myprog.exe Will run more slowly with -pg  
gprof myprog.exe gmon.out >  
profile.txt  
less profile.txt
```

Parses the output and assembles into something readable..

Summary

- For the fastest possible serial code:
 - Choose the best algorithm.
 - First, examine your within-core performance.
 - Never guess. Do experiment. Always use a profiler.
 - Look at your use of the memory hierarchy.
 - Don't revisit memory more frequently than you need to.
 - To get the most out of modern processors, you will need code that will vectorise well.
- Appendix: Cache thrashing and tiled loops.

Inner Product


c		a		b		
1x + 2y + 3z		1	2	3		x
4x + 5y + 6z	=	4	5	6	x	y
7x + 8y + 9z		7	8	9		z

N=3

```
/*  
** matrix vector multiply--a classic example  
** from: http://en.wikipedia.org/wiki/Loop\_tiling  
*/  
for (i = 0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        c[i] = c[i] + a[i][j] * b[j];  
    }  
}
```


But what if N is large?

Many Cache Misses



End up re-loading b and c many times

Fewer Cache Misses



Work through arrays in block-size: STEP

'Blocked' Loop: Code

```
/*  
** tiling of STEPxSTEP blocks  
*/  
for (i=0; i<N; i+=STEP) {  
    for (j=0; j<N; j+=STEP) {  
        for (x=i; x<min(i+STEP, N); x++) {  
            for (y =j; y<min(j+STEP, N); y++) {  
                r[x] = r[x] + p[x][y] * q[y];  
            }  
        }  
    }  
}
```

'Blocked' Loop: Run-time..

gcc v4.4.6 on BCp1

Naive nested loops	
Elapsed time:	2.585584 (s)
Tiled loop	
Elapsed time:	1.743811 (s)

So all the extra coding was worth it!

Naive nested loops	
Elapsed time:	1.900453 (s)
Tiled loop	
Elapsed time:	1.993969 (s)

icc v12.1 on BCp1

Not if your compiler will block loops for you!