



# WHY VECTORIZE?

Jim Cownie

# Overview

## Fundamental trends

- Feature sizes
- Performance, memory bandwidth
- Energy

## What they mean for architecture

## Vectors

- History
- Now

## Software

- What do I have to do to use them?

# “It's tough to make predictions, especially about the future” So we look at the past...

## Performance


FP improvement higher than integer

Still rising, but not on trend

Why not?

### Notes:

- Data from Stanford CPUdb
- This starts later than the next few graphs (no SPEC numbers before this!)
- Y-axis for these “Scaling” graphs is
  - Logarithmic
  - Ratio of metric/min(metric)




# What about other metrics?

## Clock speed

Ouch!

Hardly any change since  
2004

Performance improvements  
have to come from  
somewhere else




# What about other metrics?

## Number of Transistors

### Moore's Law:

Available transistors double every  
(1½years, 2years, 2½years)




# What about other metrics?

## Number of Cores

Didn't increase until  
after 2004

Has increased rapidly  
since

That's where lots of the  
transistors are going



# What about other metrics?

## Power

Increasing until ~2004


~Static or decreasing since

In some ways that's good

- Don't want to burn a hole in your pocket
- Energy efficiency has to improve

$$\text{Power} = CV^2f + VI_{\text{leakage}}$$

Power limit bounds frequency  
once Dennard scaling stops



# How do we use all the transistors?

Eat other system components

- Graphics
- Memory i/f
- PCI i/f

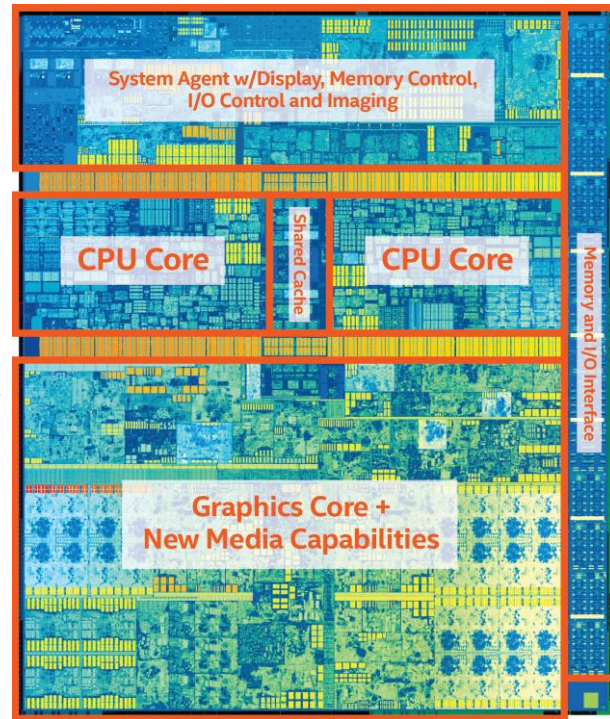
Add cache

Replicate cores

This is a laptop part.

- Two cores x two HW threads
- **256 bit (8 single or 4 double) SIMD FP unit.**
- **FMA (fused multiply add)**

Intel® 7<sup>th</sup> generation  
Core processor (“Kaby  
Lake”) released January  
2017




Data and thread parallelism are mandatory to extract all available performance.



# How do we use all the transistors for HPC?

## Intel® Xeon Phi™ 7200 series processor

- Up to 72 cores
- **Two 512 bit SIMD FPU** (16SP, 8DP, with FMA) per core
- **Vector Peak Perf:** 3+TF DP and 6+TF SP Flops
- 14nm process
- Boots normal OSES runs X86\_64 code
- Fortran, C, C++, OpenMP\*, MPI, Python,...
- Slower clock than mainstream Xeons
- 16GiB in package high BW MCDRAM



Data and thread parallelism are even more important here!

# Fundamentals: Energy

ALU ops themselves are cheap

Locality (even on die) is important, and becomes critical in the future

- Register files and caches matter (must re-use data)
- Instruction decode matters (so need to execute fewer)

Nvidia* energy numbers	2010 40nm	2017 10nm high freq	2017/2010
DP Fused Multiply Add	50 pJ	8.7 pJ	17%
3x64bit read, 1x64bit write in 8K SRAM	56 pJ	9.6 pJ	17%
On die wire energy (256b moving 10mm)	310 pJ	200 pJ	<b>65%</b>
DRAM Access	10,000 pJ	1,700 pJ	17%

Vectors and vector registers help both issues


# Memory Bandwidth keeps getting (relatively) worse

Smaller balance is better (it's flops/bw)

Lots of variation within a year

But the **best** machine now is  
>10x worse than in 1997

Now need > 10 ops per load or store to achieve peak FLOPS



# History

## Ancient history

- Illiac IV (1972)
- CDC Star (1974) – Memory to memory architecture
- Cray 1 – Vector registers

Lots of memory bandwidth/flop !

Main benefits come from good pipelining of operations

## More recent history

### X86

- SSE(1999), SSE2(2001), SSE3(2004), SSE4(2007) 8x128b XMM registers
- AVX1(2011), AVX2(2013) 16x256b YMM registers
- Intel® Xeon Phi™ coprocessor (2012) 16x512b ZMM registers, 2017: AVX512

ARM : Neon, Power: AltiVec, ...

# Achieving Performance

Every code has a performance limiting resource

- Otherwise it would run faster until it did!

We need to understand what the limiting resource is

- Architect: so we can provide more of it
- Programmer: so we can rewrite our code to avoid it

The commonest limits for HPC codes are

- Memory bandwidth
- Floating point resources

Vectors help

- Efficient way to drive lots of ALUs
- Encourage efficient use of memory bandwidth

# The Fundamental Idea

Expose data parallelism to the hardware

If it knows, it can exploit it, but the compiler has to tell it...

```
void xpy (float *x, float const * y, int len) {  
    for (int i=0; i<len; i++)  
        x[i] = x[i] + y[i];  
}
```

We seem to be operating on a whole vector, but it's not obvious...

Can the compiler see it with this code?

No: src and dest can alias; need restrict qualifiers

# Why can't that code just vectorize?

Consider this invocation

```
float x[10];  
for (int i=0; i<10; i++)  
    x[i] = (float)i;  
xpy(&x[1],&x[0], 9);
```

```
void xpy(float *x,float const *y,int len)  
{  
    for (int i=0; i<len; i++)  
        x[i] = x[i] + y[i];  
}
```

Correct result is 0,1,3,6,10,15,21,28,36,45

If you force vectorization (with SIMD width of 4) you get  
0,1,3,6,10,**9,11,13,15,17**

The compiler cannot vectorize unless you give it more information.

`restrict` on a pointer says “it does not alias” => no order between accesses via it and other pointers

# Array of Structures and Structure of Arrays

C/C++ make it natural to write code in array of structures (AOS) form like this :-

```
void xpy (element *x, element const * y, int len) {  
    for (int i=0; i<len; i++)  
        x[i].field = x[i].field + y[i].field;  
}
```



Inefficient; requires scatter/gather and wastes memory bandwidth

Better to rework to structure of arrays form (but a hard change in a large codebase ☹)



# What does the compiler need to know?

## Aliasing information

- Use `restrict` (or Fortran which has always specified no overlap of arguments)
- Use `omp simd` or other pragmas (there are way too many!)

## Alignment information `omp simd aligned`

## Expected iteration counts

- Pragmas for these too

## Look at vectorization reports `-vec-report n`

## When the hints fail... force the issue

- Use OpenMP 4.5 `simd` directives  
    `#pragma omp simd (!$omp simd in Fortran)`
- Can annotate vector callable functions
- Can be used without also using OpenMP parallelism (to mix with Intel® Threading Building Blocks, for instance) `-fopenmp-simd`

# Low level interface: Intrinsics

Another way of writing something very close to assembler code.

Like this

```
register __m128d sigma2;  
sigma2 = _mm_setzero_pd()  
v_8_9 = _mm_load_pd(&Matrix[0][8]);  
xsigma1 = _mm_add_pd(_mm_mul_pd(v_8_9, v_8_9), xsigma1);
```

Explicitly operating on SIMD-width vectors

Slightly higher than machine code, but not much! (note the explicit 128b width!)

# Low level interface

## Plus

- Highest performance is possible
- You control the horizontals and the verticals

## Minus

- Like all ASM code it's hard work writing this stuff
- Not portable
  - Certainly not to non-X86
  - May not be highest performance on future cores (e.g. SSE to AVX2 or AVX2 to KNC or AVX512)

Use only if you really need to (I try to avoid them)  
Read the compiler's asm output before you start.

# Why Vectorize rather than Parallelize?

Parallelization redistributes work to get you an answer sooner

- Using four cores I may be able to execute 4 CPU seconds of work in ~1 elapsed second (but I still consumed 4 CPU seconds of resource)

Vectorization reduces the total amount of work

- Using vectorization I may be able to compute 4 times faster on a single CPU, reducing that 4 CPU seconds to 1 CPU second

If you're charged for machine time vectorization should save you money, whereas parallelization won't!

# Memory Bandwidth

Often the limiting factor in HPC codes, rather than raw FLOPS

Intel® Advisor can help

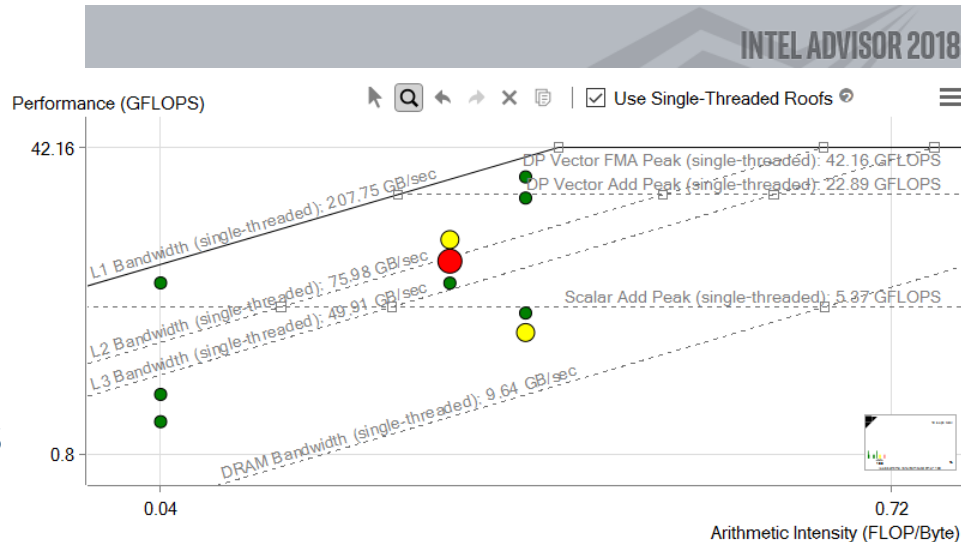
- Roofline model
- Vectorization reports integrated back into code

# Find Effective Optimization Strategies

Intel Advisor: Cache-aware roofline analysis

## Roofline Performance Insights

- Highlights poor performing loops
- Shows performance “headroom” for each loop
  - Which can be improved
  - Which are worth improving
- Shows likely causes of bottlenecks
- Suggests next optimization steps



# Find Effective Optimization Strategies

Intel Advisor: Cache-aware roofline analysis

## Roofs Show Platform Limits


- Memory, cache & compute limits

## Dots Are Loops

- Bigger, red dots take more time so optimization has a bigger impact
- Dots farther from a roof have more room for improvement

## Higher Dot = Higher GFLOPs/sec

- Optimization moves dots up
- Algorithmic changes move dots horizontally



## Which loops should we optimize?

- A and G are the best candidates
- B has room to improve, but will have less impact
- E, C, D, and H are poor candidates

[Roofline tutorial video](#)

# Conclusions

Vectors are here to stay

Code needs to be designed with that in mind

- SOA instead of AOS
- Give the compiler all the information you can

On Intel® Xeon Phi™ processor and AVX-512 enabled Xeons you're only using 1/16<sup>th</sup> of the SP float capability if you haven't vectorized

Even on a laptop Intel® processors you're down by 8x today!

Good Luck!



# References

Stanford CPUDB <http://cpudb.stanford.edu/>

Energy numbers from Nvidia

<http://research.nvidia.com/publication/gpus-and-future-parallel-computing>  
page 9

Stream <http://www.cs.virginia.edu/stream/>

# Free Software

As students at a degree awarding institution you qualify:-

<https://software.intel.com/en-us/qualify-for-free-software/student>

Free download of Intel® Parallel Studio XE

**BACKUP**

# Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2016, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.


Notice revision #20110804

# INTEL® COMPILER PERFORMANCE

- Compare SpecOMP12 on identical hardware
- Single node in a Cray\* XC30: 2xIntel® Xeon E5-2697 v2, Aug 2017 data collected by Indiana University
- It is hard to find such results; most vendor runs just use Intel® compilers
  - I wonder why? 😊
  - I couldn't find SpecINT/FP runs to compare!

>20% performance improvement over GNU compilers!

## SpecOmp12 Performance by Compiler



Data from <http://spec.org> downloaded on 8 September 2017 ("omp20212-results-20170907-095222")

### Optimization Notice

Copyright © 2017, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.



