

OpenMP 2

Multi-threaded Programming



theory




practice

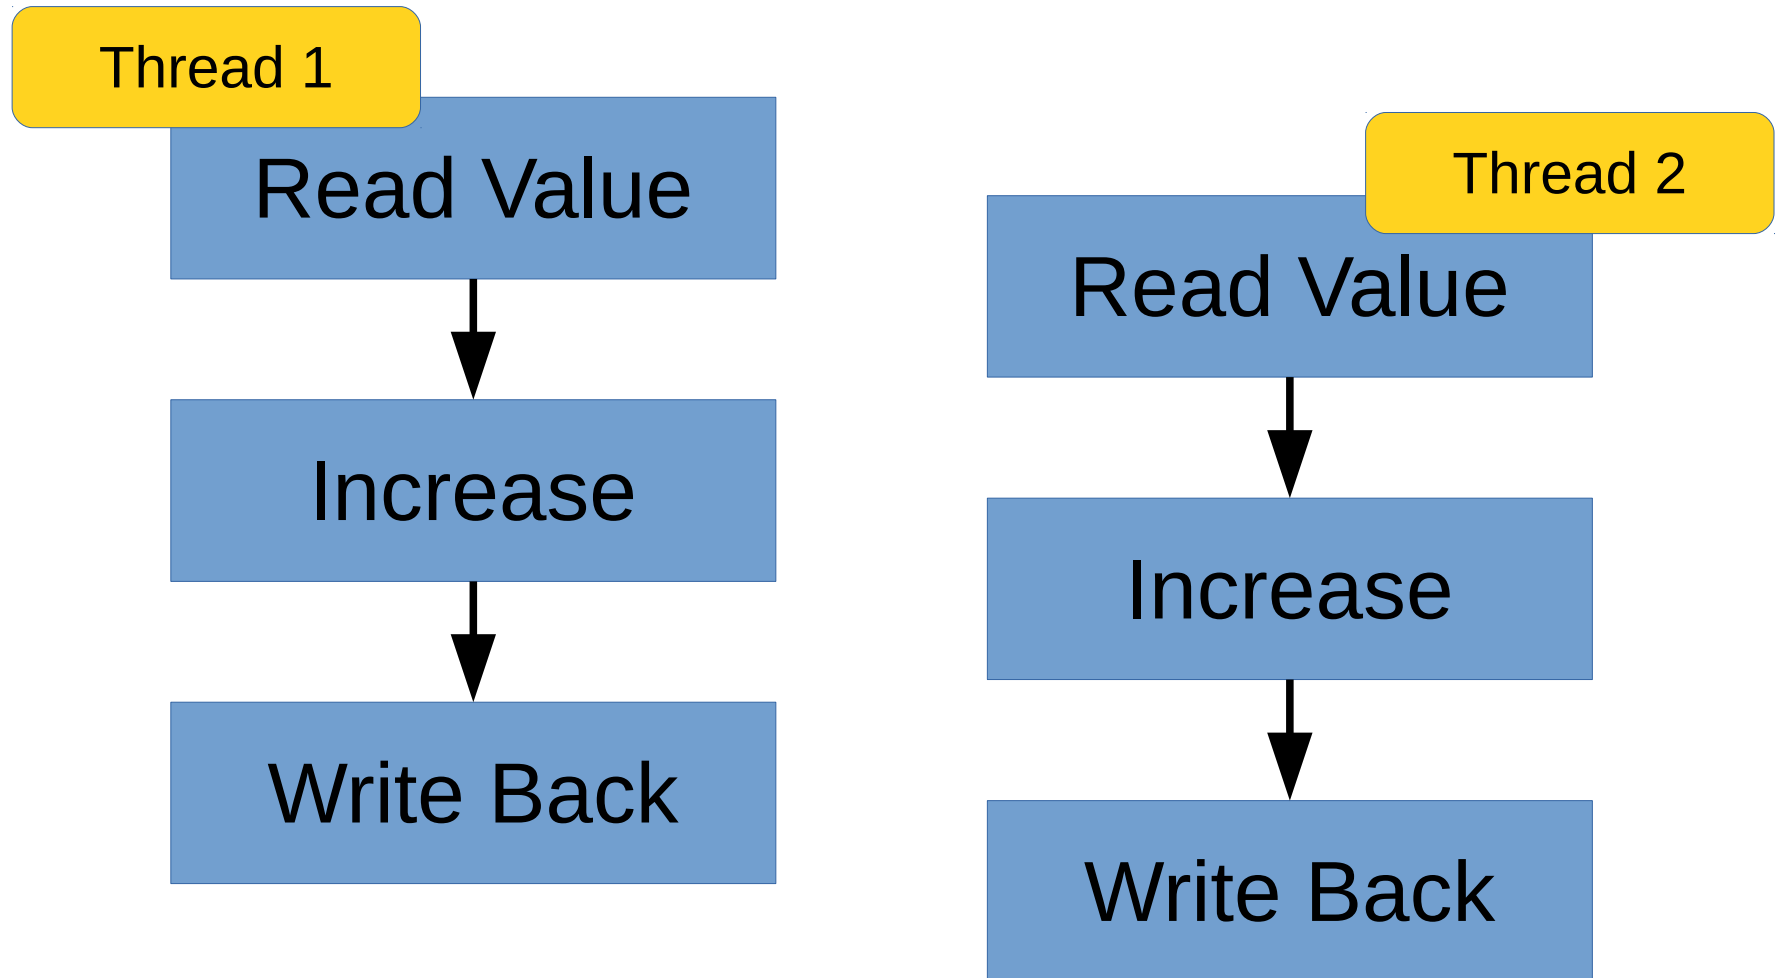
Overview

- Parallel programming – new opportunities! But also new bugs and new performance gotchas:
 - Race condition bugs, barriers and load balance.
- We'll consider several versions of a numerical integration example and review performance.
- Avoiding unwanted barriers.

A Serial Code Sequence



A Threaded Code Sequence



Parallel Bugs: Race Conditions

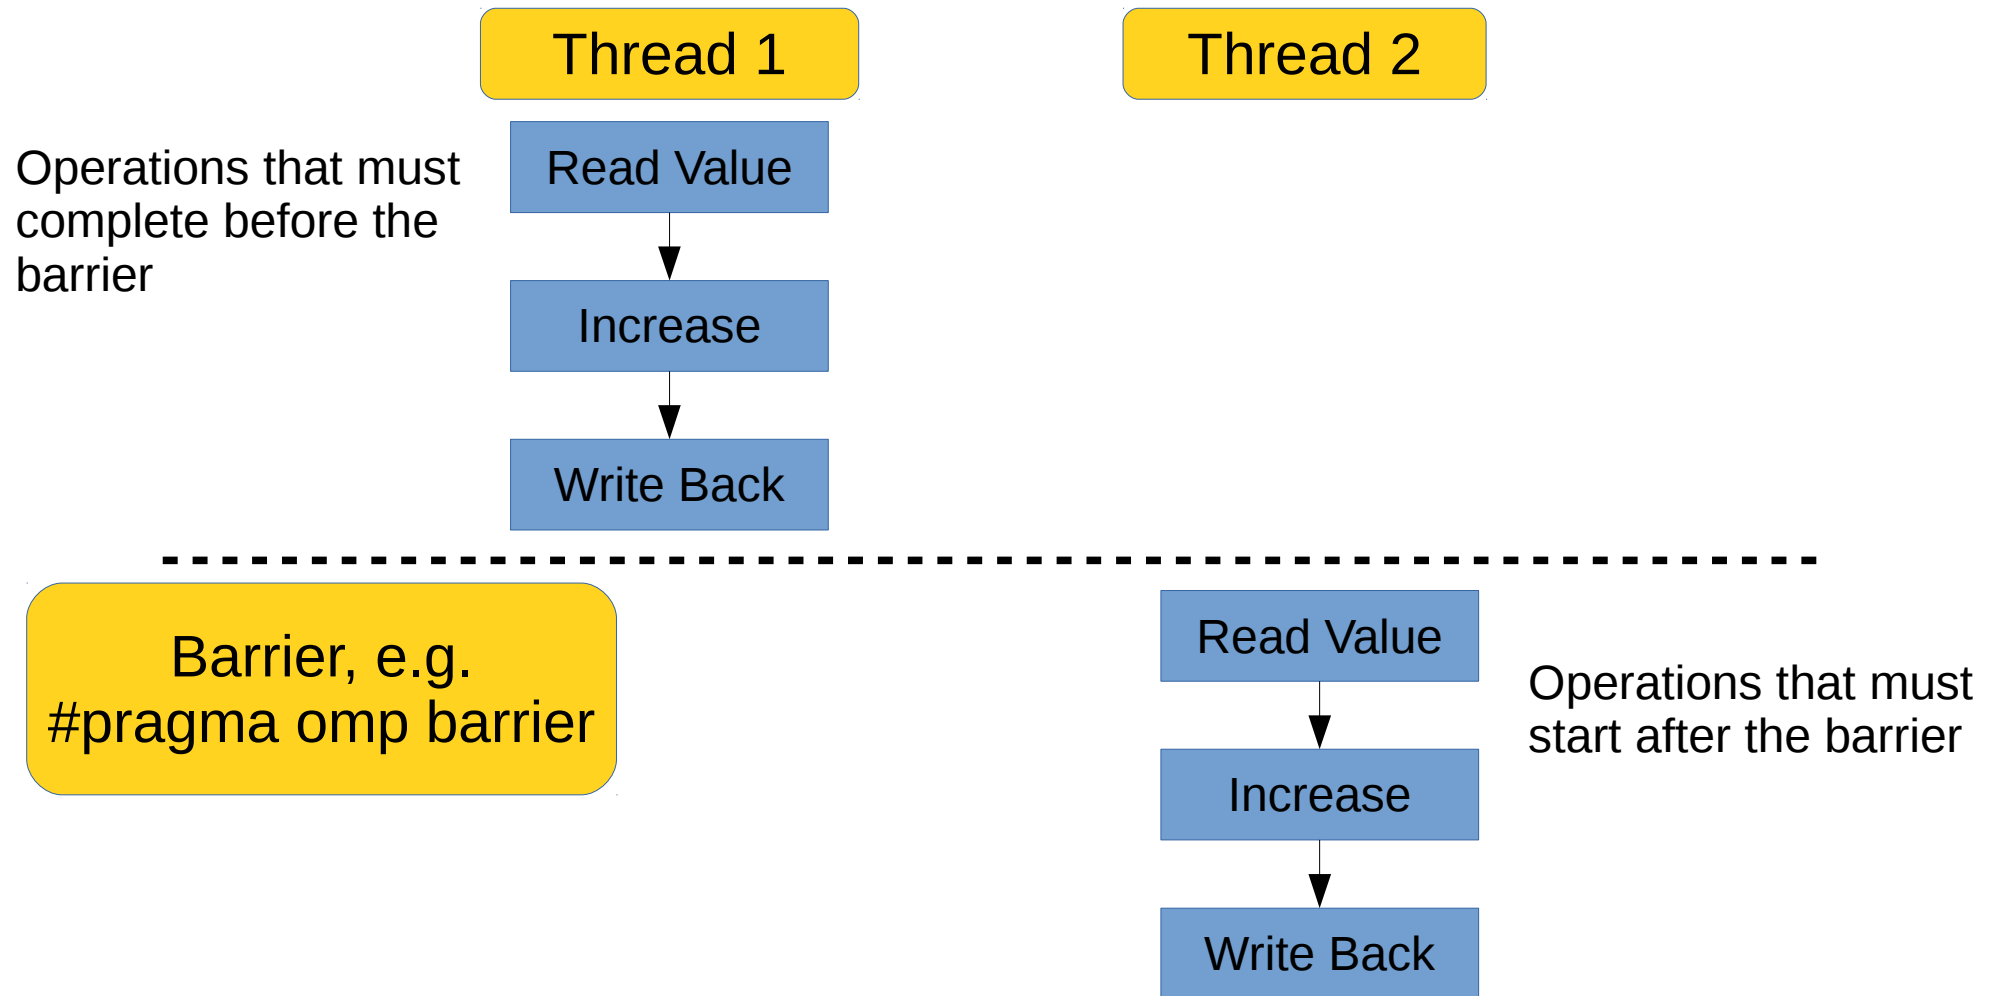
Thread 1	Thread 2		Integer Value
			0
Read value		←	0
Increase value			0
Write back		→	1
	Read value	←	1
	Increase value		1
	Write back	→	2

Intended

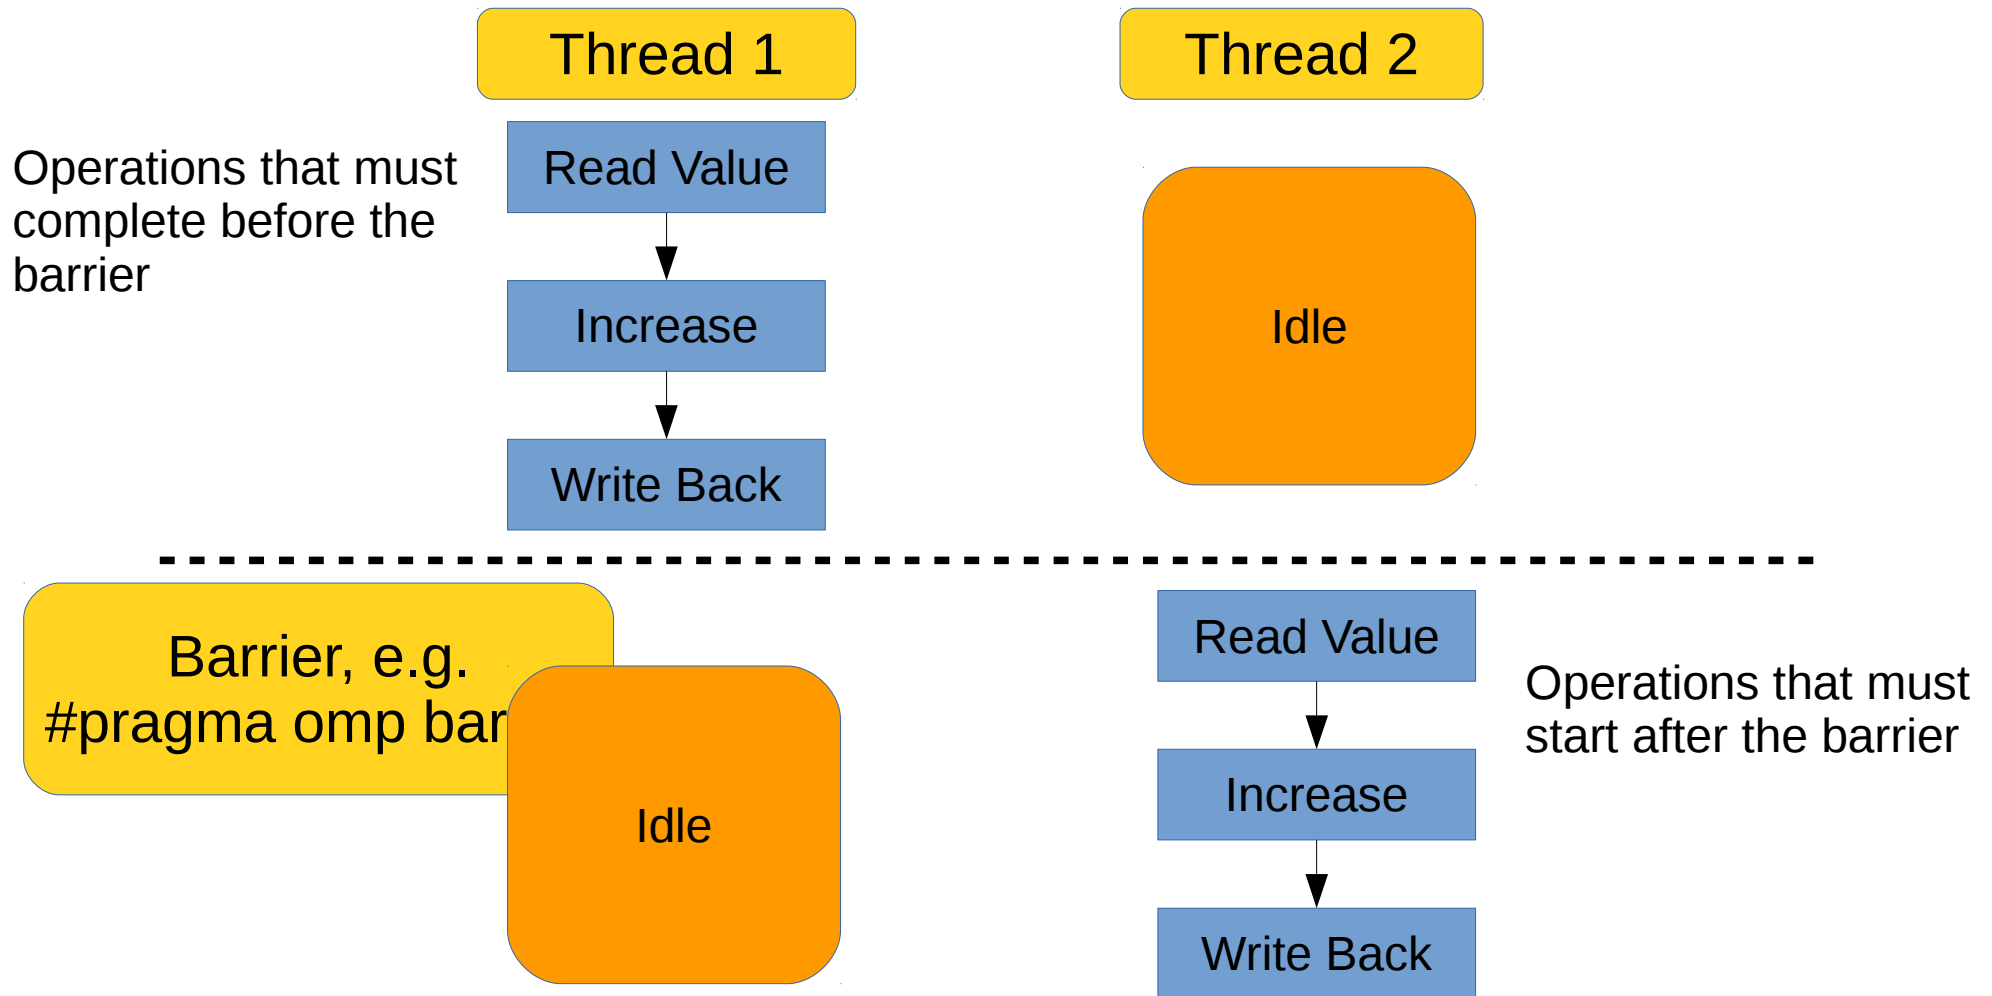
Thread 1	Thread 2		Integer Value
			0
Read value		←	0
	Read Value	←	0
Increase value			0
	Increase value		0
Write back		→	1
	Write back	→	1

What we got


Introducing a Barrier



Introducing a Barrier




Performance: Synchronisation & Load Balance



- We might use a barrier to avoid a '**race condition**' error. (e.g. https://en.wikipedia.org/wiki/Race_condition)
- **Parallel for loops have an implicit terminating barrier.**
- **A critical region or 'mutex' (perhaps to compute a total) will force a synchronisation.**

Example: numerical integration



(e.g. estimate pi by numerically integrating $4/(1+x^2)$)

Shared Accumulator


```
#pragma omp parallel for shared(sum) private(x)
  for (ii=0; ii<num_steps; ii++) {
    x = (ii+0.5)*step;
#pragma omp critical
    sum += 4.0/(1.0+x*x); /* all threads access! */
  }

/* master thread only */
pi = step * sum;

printf("pi is:\t\t%.16f\n", pi);
printf("error is:\t%.16f\n", fabs(pi - PI25DT));
```

Syntactically fine, but what's the problem?

Shared Accumulator

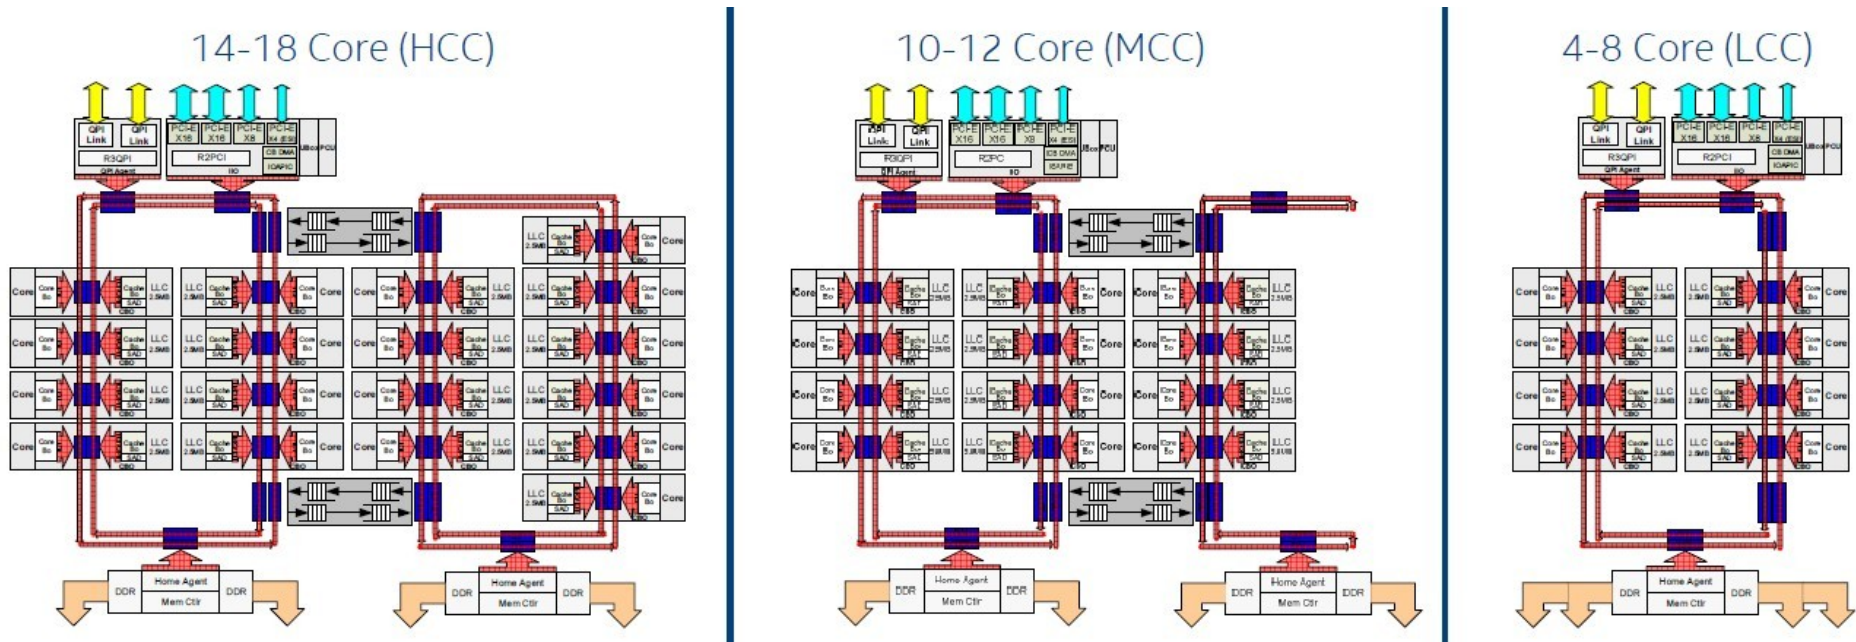


Possible Reasons for Load Imbalance?

- Unequal division of labour.
- Independent clocking of cores (to remain inside thermal limits). For example, interaction between AVX instructions and turbo frequencies,
 - see e.g.:
<http://www.anandtech.com/show/8423/intel-xeon-e5-version-3-up-to-18-haswell-ep-cores-/5>

Possible Reasons for Load Imbalance?

- NUMA: Current processors use bus architectures to access memory, and they are less symmetric than you might expect!



Performance: Schedule

- **Static**

- Iterations are divided evenly among threads. Extra control via chunk size. Chunks assigned on a 'round-robin' basis.

- **Dynamic**

- Threads will come back for more once they have completed their block—faster threads free to acquire more work.

- **Guided**

- As per dynamic but with initial large blocks (can save some overheads).

Performance: Schedule

```
#define CHUNKSIZE 10
#define N 100
int main (int argc, char *argv[]) {
    chunk = CHUNKSIZE;
#pragma omp parallel shared(a,chunk) \
                    private(i,tid,nthreads)
    {
        ...
#pragma omp for schedule(static,chunk)
        for (i=0; i<N; i++)
        {
            a[i] = i * 1.0;
            printf("Thread %d: a[%d]= %f\n",tid,i,a[i]);
        }
        printf("Thread %d done.\n",tid);
    } /* end of parallel section */
}
```

The 'nowait' Clause

```
#pragma omp parallel
{
    #pragma omp for nowait
        for (ii=1; ii<N; ii++)
            a[ii] = (a[ii] + a[ii-1]) /2.0;
    #pragma omp for
        for (ii=1; ii<N; ii++)
            c[ii] = d[ii]/e[ii];
}
```

- Allows you to avoid the implicit barrier at the end of a worksharing for loop.
- Only possible if your loops are independent (note different data structures).

More on Load (Im)Balance

- Consider that you have 100 cores at your disposal:
 - Uncommon even a few years ago, but now close to routine..
 - Each core will do 1% of the work (given a task which we can parallelise well).
 - What if one processor was running just over 1% slower than the rest?

More on Load (Im)Balance

- Consider that you have 100 cores at your disposal:
 - Uncommon even a few years ago, but now close to routine..
 - Each core will do 1% of the work (given a task which we can parallelise well).
 - What if one processor was running just over 1% slower than the rest?

You'd be better off without it!

More on Load (Im)Balance

- This example gives us a feel for how important load balance can be, as core counts rise..
- Happily load tends to be balanced in (data parallel) work-sharing loops.
- But this is not the case for task parallel patterns.

Summary

- Use OpenMP wisely:
 - Concurrent threads gives us lots of opportunities
 - But also introduces new potential pitfalls
 - Race Condition Bugs
 - Load Imbalance
 - And methods to combat:
 - Good design
 - Schedule
 - nowait