

# OpenMP 3

## Multi-threaded Programming



theory




practice

# Overview

- Continuing with our numerical integration example and reviewing performance.
- We now encounter ‘false sharing’ as a performance gotcha.
- Also consider how many times you create and destroy teams of threads, because both operations carry an overhead.
- Let’s be clear about what we can hope to gain by going parallel.

# Example: numerical integration



(e.g. estimate pi by numerically integrating  $4/(1+x^2)$ )

# Shared Accumulator

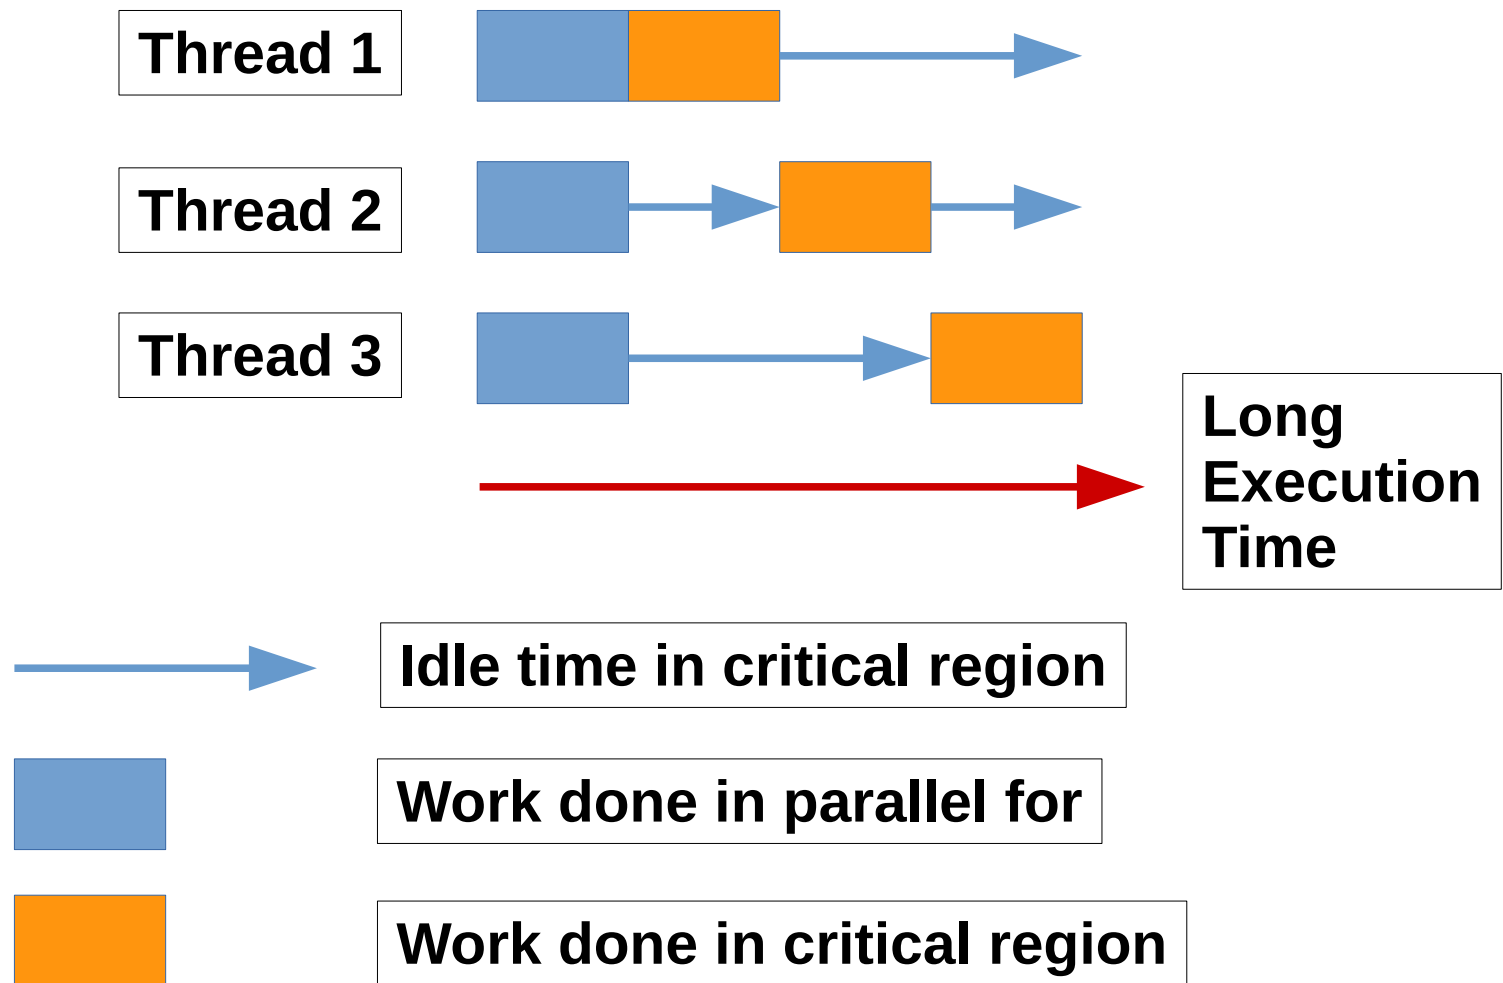
```
#pragma omp parallel for shared(sum) private(x)
  for (ii=0; ii<num_steps; ii++) {
    x = (ii+0.5)*step;
#pragma omp critical
    sum += 4.0/(1.0+x*x); /* all threads access! */
  }

/* master thread only */
pi = step * sum;

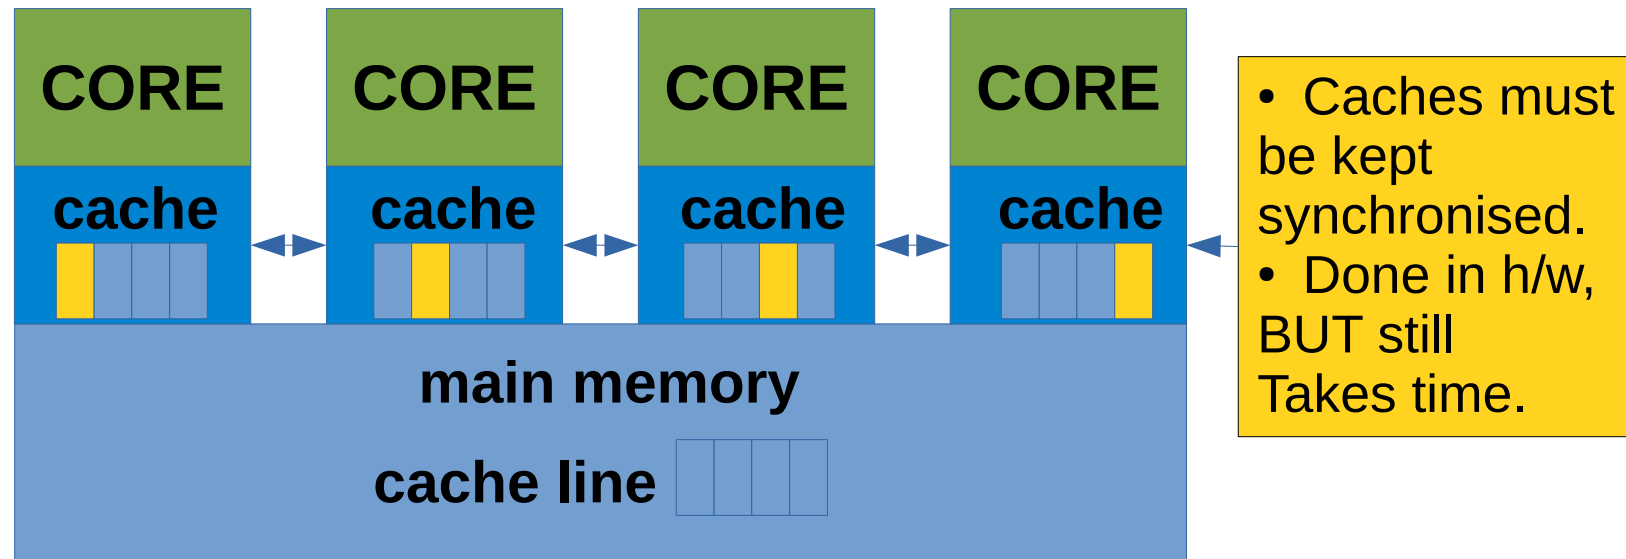
printf("pi is:\t\t%.16f\n", pi);
printf("error is:\t%.16f\n", fabs(pi - PI25DT));
```

**Syntactically fine, but load imbalance  
Gives us poor performance**

# Shared Accumulator



# Performance: Cache Coherency & False Sharing



**'False sharing'** is an example of cache thrashing in a concurrent context. It occurs when two threads access different words that are on the same cache line.

**e.g. A small array that will fit in cache, one cell used per processor.**

x86: Cache lines are 64 bytes

Explore: `/sys/devices/system/cpu/cpu0/cache`

# Array of Accumulators


```
double sum[NUM_THREADS];          /* array to store partial sums */

#pragma omp parallel shared(sum,pi) private(x,thread_id)
{
    thread_id = omp_get_thread_num();
    sum[thread_id] = 0.0;
#pragma omp for
    for (ii=0; ii<num_steps; ii++) {
        x = (ii+0.5)*step;
        sum[thread_id] += 4.0/(1.0+x*x);
        /* .. note that whole array will likely be in cache */
    }
}

/* back to just master thread */
/* total the partial sums */
for (ii=0; ii<NUM_THREADS; ii++) {
    pi += sum[ii];
}
pi *= step;
```

Nice try, but..

# Array of Accumulators



Constantly 'thrashing cache'



# Private Accumulators


```
#pragma omp parallel shared(pi) private(x) \  
    firstprivate(sum)  
{  
#pragma omp for  
    for (ii=0; ii<num_steps; ii++) {  
        x = (ii+0.5)*step;  
        /* partial sum is now private in each thread  
         * note the use of firstprivate to ensure an  
         * initial value for the accumulator */  
        sum += 4.0/(1.0+x*x);  
    }  
    /* a critical section is a sequence of mutual  
     * exclusion (mutex) locks which ensures access  
     * for only one thread at a time */  
#pragma omp critical  
    pi += sum;                    /* but total is shared */  
}
```

**Much better, but still not best..**

# Reduction


```
#pragma omp parallel for reduction(+:sum) private(x)
  for (ii=0;ii<num_steps; ii++){
    x = (ii+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
    /* ..note that sum is now in a reduction */
  }
```

- Much neater! (and potential better load balance)
- Several reduction operators available: +,-,\*,logical & bitwise ops
- Fortran also allows 'min' and 'max'




**Better load balance  
through tree structure**

# Performance: Cost of Team Creation & Destruction

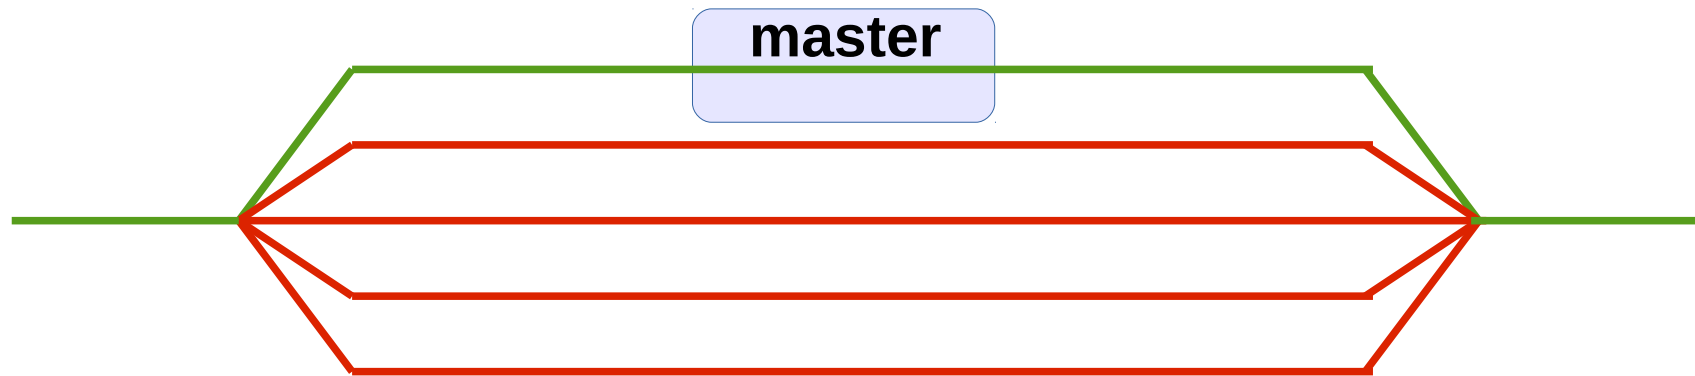


# Creating & Destroying Teams



```
#pragma omp parallel for  
...  
#pragma omp parallel for  
...
```

# Creating & Destroying Teams



```
#pragma omp parallel  
#pragma omp for  
...  
#pragma omp for  
...
```

# Threading Overheads

On my (old) (4-core, single-socket, i3) desktop..

```
for(j=0; j<REPEAT; j++) {  
#pragma omp parallel for  
  for (i=0; i<N; i++) {  
    a[i] = b[i];  
  }  
}
```

```
real 0m3.526s  
user 0m13.566s  
sys 0m0.024s
```

```
#pragma omp parallel private(j)  
  for(j=0; j<REPEAT; j++) {  
#pragma omp for  
    for (i=0; i<N; i++) {  
      a[i] = b[i];  
    }  
  }
```

```
real 0m1.090s  
user 0m3.944s  
sys 0m0.004s
```

Here we see the cost of repeatedly creating and destroying teams of threads.

# Threading Overheads

```
#pragma omp parallel for shared(sum)
private(x,x2)
for (ii=0; ii<num_steps; ii++) {
    x = (ii+0.5)*step;
    x2 = 4.0/(1.0+x*x);
#pragma omp critical
    sum += x2; /* all threads access! */
}
```

Here we see that the critical region accounts for the majority of the run time.

ON SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs usec/call	Inclusive Name
100.0	48	4:56.696	1	1	296696151 .TAU application
100.0	0.0423	4:56.647	1	1	296647246 parallel begin/end [OpenMP]
100.0	0.0535	4:56.647	1	1	296647204 parollefor (parallel begin/end) [OpenMP location: file:/panfs/panasas01/isys/ggdagw/hpc-course/openmp/examples/example2/parallel_shared_pi.c <26, 32>]
100.0	0.0357	4:56.647	1	1	296647150 for enter/exit [OpenMP]
100.0	1:37.962	4:56.647	1	2.5E+07	296647115 parollefor (loop body) [OpenMP location: file:/panfs/panasas01/isys/ggdagw/hpc-course/openmp/examples/example2/parallel_shared_pi.c <26, 32>]
64.9	35,857	3:12.667	2.5E+07	2.5E+07	8 critical enter/exit [THROTTLED]
52.9	1:58.603	2:36.809	2.5E+07	2.5E+07	6 critical (critical enter/exit) [OpenMP location: file:/panfs/panasas01/isys/ggdagw/hpc-course/openmp/examples/example2/parallel_shared_pi.c <30, 31>] [THROTTLED]

# Potential Benefits: Amdahl's Law

$$speedup = \frac{T_1}{T_p}$$

Best you can ever get

$$speedup_{max} = \frac{1}{1-P}$$

$$speedup_{ideal} = \frac{1}{\frac{P}{N} + S}$$

If you've coded it perfectly

	Ideal <i>parallel</i> speed-up		
N	P=0.5	P=0.9	P=0.99
10	1.82	5.26	9.17
100	1.98	9.17	50.25
1000	1.99	9.91	90.99
10000	1.99	9.91	99.02
100000	1.99	9.99	99.90

- $T_1$  is the execution time on a single processor
- $T_p$  is the execution time on a parallel computer
- $S$  is the serial fraction of the program
- $P$  is the parallelisable fraction of the program
- $N$  is the number of processors available

**'strong' scaling**

Amdahl, Gene (1967). "Validity of the Single Processor Approach to Achieving Large-Scale Computing Capabilities". AFIPS Conference Proceedings (30): 483-485



# Potential Benefits: Gustafson's Law

'weak' scaling

$$\text{speedup} = \frac{T_1}{T_p} = \frac{T_a(x) + N \times T_b(x)}{T_a(x) + T_b(x)} \rightarrow N$$

- $x$  a measure of problem size
- $N$  processors
- $T_a(x)$  fraction of time spent executing the serial part of the program
- $T_b(x)$  fraction of time spent executing the parallel part

As  $x$  increases:

- $T_a \rightarrow 0$
- $T_b \rightarrow 1$

It is assumed that  $T_a$  is the minor fraction. Its relative importance will diminish with an increase in  $x$  and ultimately, speed-up will approach  $N$ .

“Reevaluating Amdahl's Law”, John L. Gustafson, Communications of the ACM 31(5), 1988. pp. 532-533.

# Summary

- Use OpenMP wisely:
  - Watch out for load imbalance and false sharing.
  - Use barriers (only) when you need them.
  - Don't create and destroy more thread teams than you need to.
  - Explore the use of 'schedule' if you suspect asymmetries and imbalance.
  - Be savvy about what extra performance more threads can bring you.