

Libraries



Rationale



Newton: "If I have seen further it is by standing on the shoulders of giants"

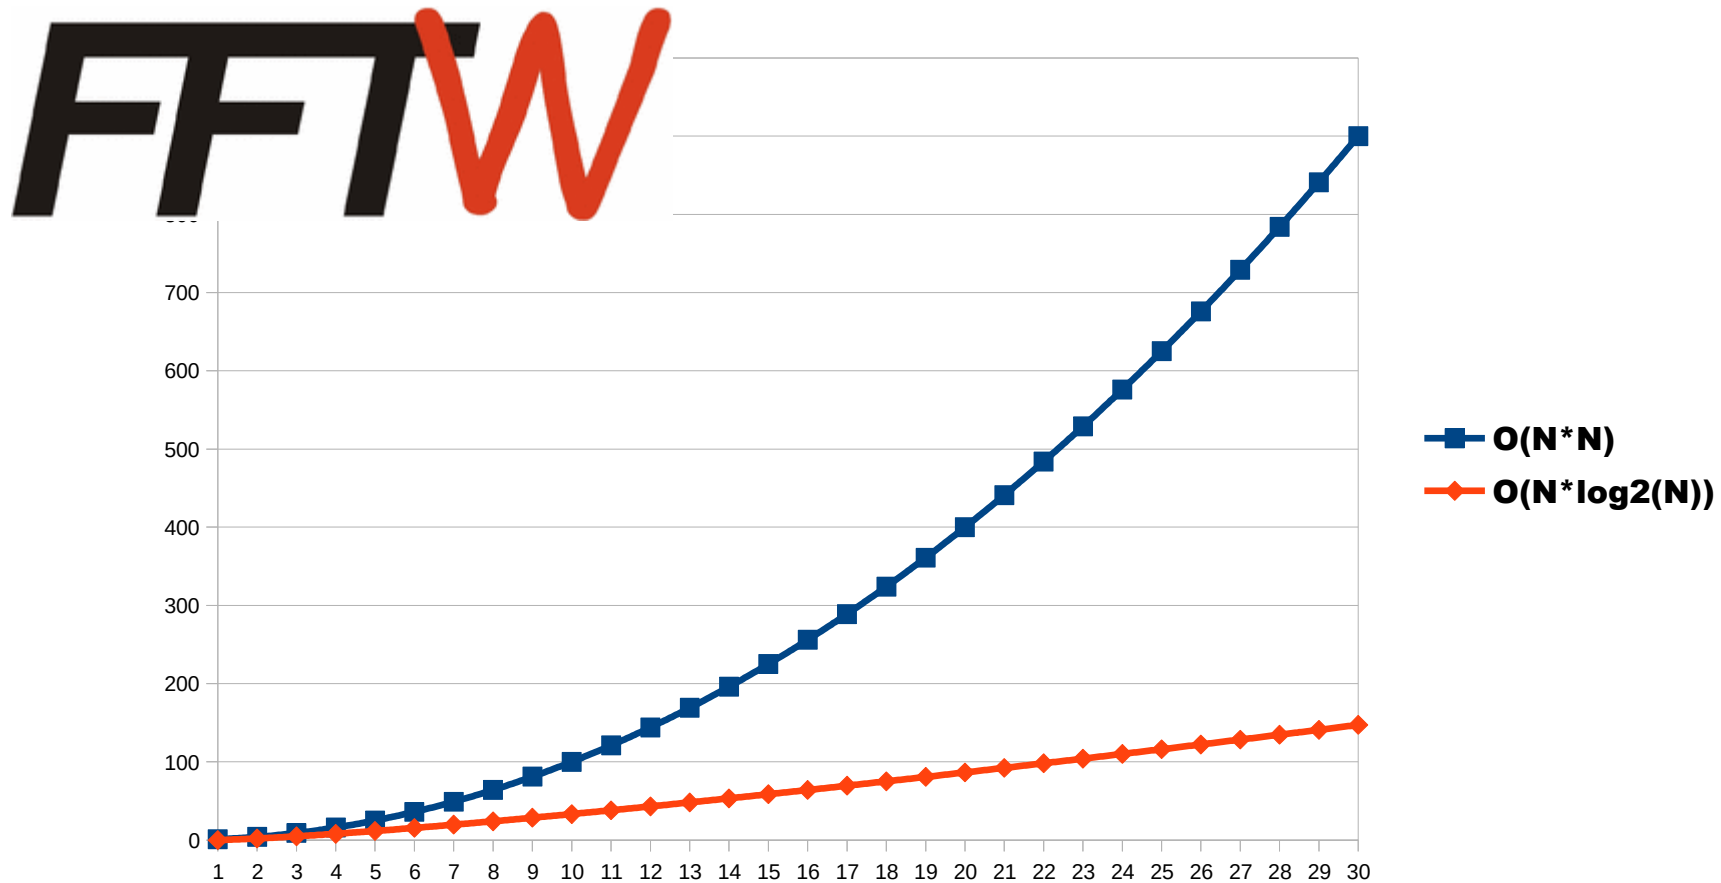
- Saves time
- Less buggy
- More efficient implementation
- Access to parallel algorithms
- But...

The only way to appreciate an algorithm and the challenges of implementing it in parallel is to code it.

Libraries and the 7 Dwarfs

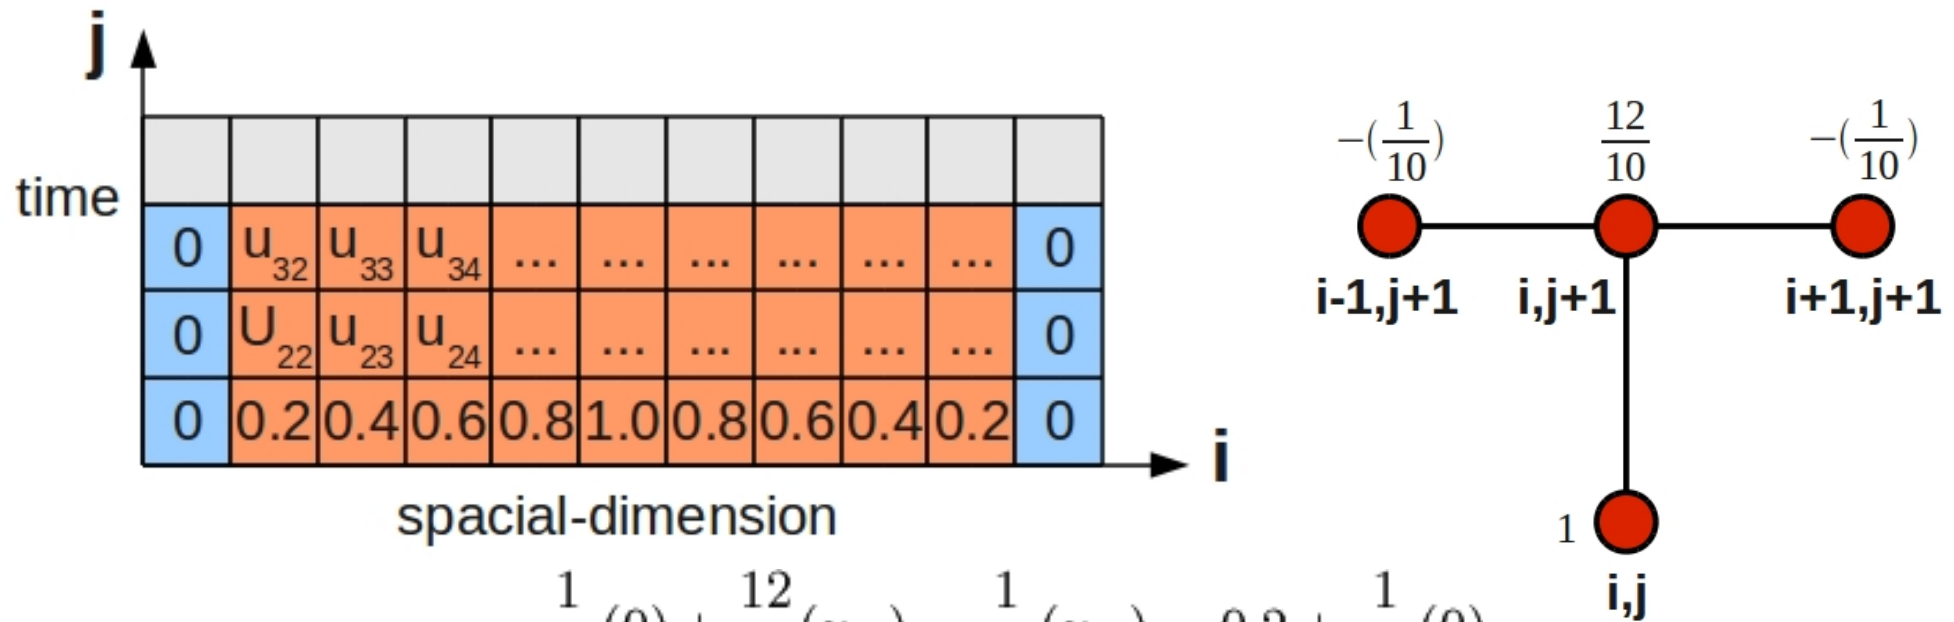
Computational kernel	Example Library/Package
Spectral	FFTW
Dense linear algebra	BLAS + LAPACK, ATLAS, ScaLAPACK, PLASMA, MAGMA, Intel MKL
Sparse linear algebra	PetSc, Trilinos
N-Body	GRAPE, Fast Multipole Method, OpenMM
Structured grids	CHOMBO, OPS
Unstructured grids	ELMER (FEM), OP2
Map Reduce	Hadoop

Spectral: $N \log(N)$ is worth having!



A Fourier transform is $O(N^2)$ operations. An FFT is $O(N \log(N))$.

Dense Linear Algebra: Solving PDEs



$$\frac{1}{10}(0) + \frac{12}{10}(u_{22}) - \frac{1}{10}(u_{23}) = 0.2 + \frac{1}{10}(0)$$

$$\frac{1}{10}(u_{22}) + \frac{12}{10}(u_{23}) - \frac{1}{10}(u_{24}) = 0.4$$

$$\begin{bmatrix} 1.2 & -0.1 & 0 & 0 & \dots & 0 \\ -0.1 & 1.2 & -0.1 & 0 & \dots & 0 \\ \vdots & & & & & \vdots \\ \dots & 0 & -0.1 & 1.2 \end{bmatrix} \times \begin{bmatrix} u_{22} \\ u_{23} \\ u_{24} \\ \vdots \end{bmatrix} = \begin{bmatrix} 0.2 + 0.1(0) \\ 0.4 \\ \vdots \\ 0.2 + 0.1(0) \end{bmatrix}$$

Dense Linear algebra: BLAS

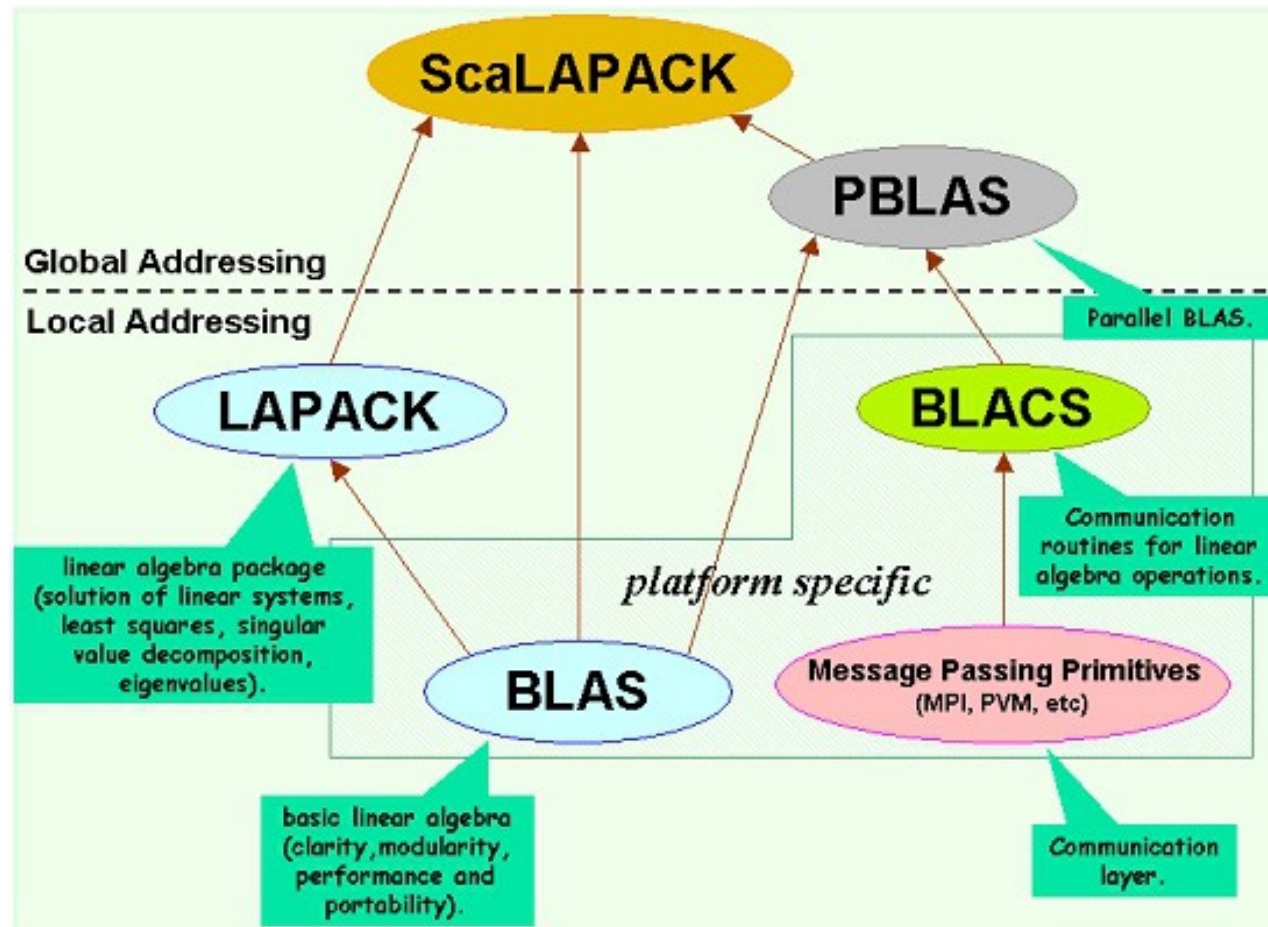
complexity of routine	# of floating point operations	# of memory accesses	Ratio of memory access to flops (lower is better)
BLAS-1 (vector-vector)	$2n$	$3n$	$3/2$
BLAS-2 (matrix-vector)	$2n^2 + n$	$n^2 + 3n$	$\sim 1/2$
BLAS-3 (matrix-matrix)	$2n^3$	$3n^2$	$3/(2n)$

- $3/(2n)$ is small for large n
- We know that memory accesses are expensive and limit performance
- BLAS-3 routines have the best chance of approaching theoretical peak performance of the CPU

Automatic Tuning of Linear Algebra Software (ATLAS)

- 'Blocked' algorithms are required to minimise the number of cache misses and hence gain best performance.
- The best block size is dependent upon the cache details for a particular microprocessor.
- ATLAS tries a number of experiments to empirically (and automatically) determine the best values for certain parameters, such as block size, length of loop to unroll etc.

ScaLAPACK: MPI Parallel



Fast vendor BLAS libraries

- Intel's Math Kernel Library (MKL)
- Nvidia's CuBLAS
- AMD's ACML
- Open source clBLAS (on github, clFFT too!)

Multi- and Many-core Versions



- <http://icl.cs.utk.edu/plasma>
- <http://icl.cs.utk.edu/magma>



TEXAS ADVANCED COMPUTING CENTER

Powering Discoveries That Change The World

<https://www.tacc.utexas.edu/research-development/tacc-software/gotoblas2>

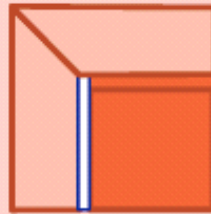


A New Generation of Software:

Parallel Linear Algebra Software for Multicore Architectures (PLASMA)

Software/Algorithms follow hardware evolution in time

LINPACK (70's)
(Vector operations)



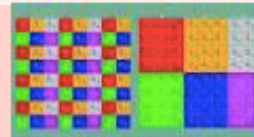
Rely on
- Level-1 BLAS operations

LAPACK (80's)
(Blocking, cache friendly)



Rely on
- Level-3 BLAS operations

ScaLAPACK (90's)
(Distributed Memory)



Rely on
- PBLAS Mess Passing

PLASMA (00's)
New Algorithms
(many-core friendly)



Rely on
- a DAG/scheduler
- block data layout
- some extra kernels

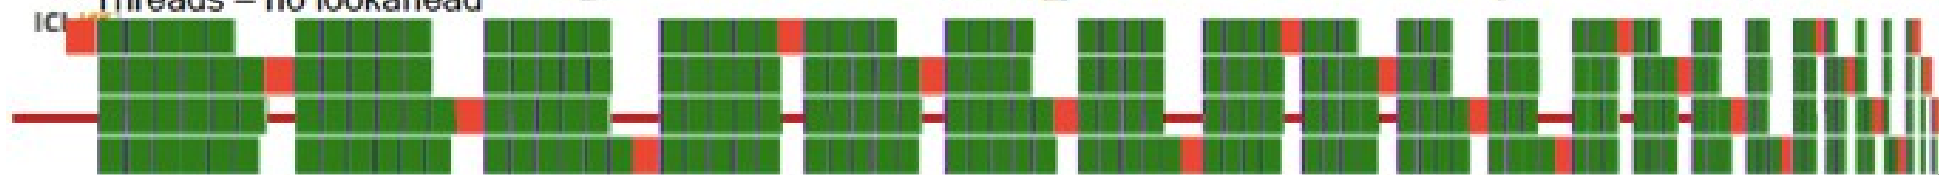
Those new algorithms

- have a very **low granularity**, they scale very well (multicore, petascale computing, ...)
- **removes a lots of dependencies** among the tasks, (multicore, distributed computing)
- **avoid latency** (distributed computing, out-of-core)
- **rely on fast kernels**

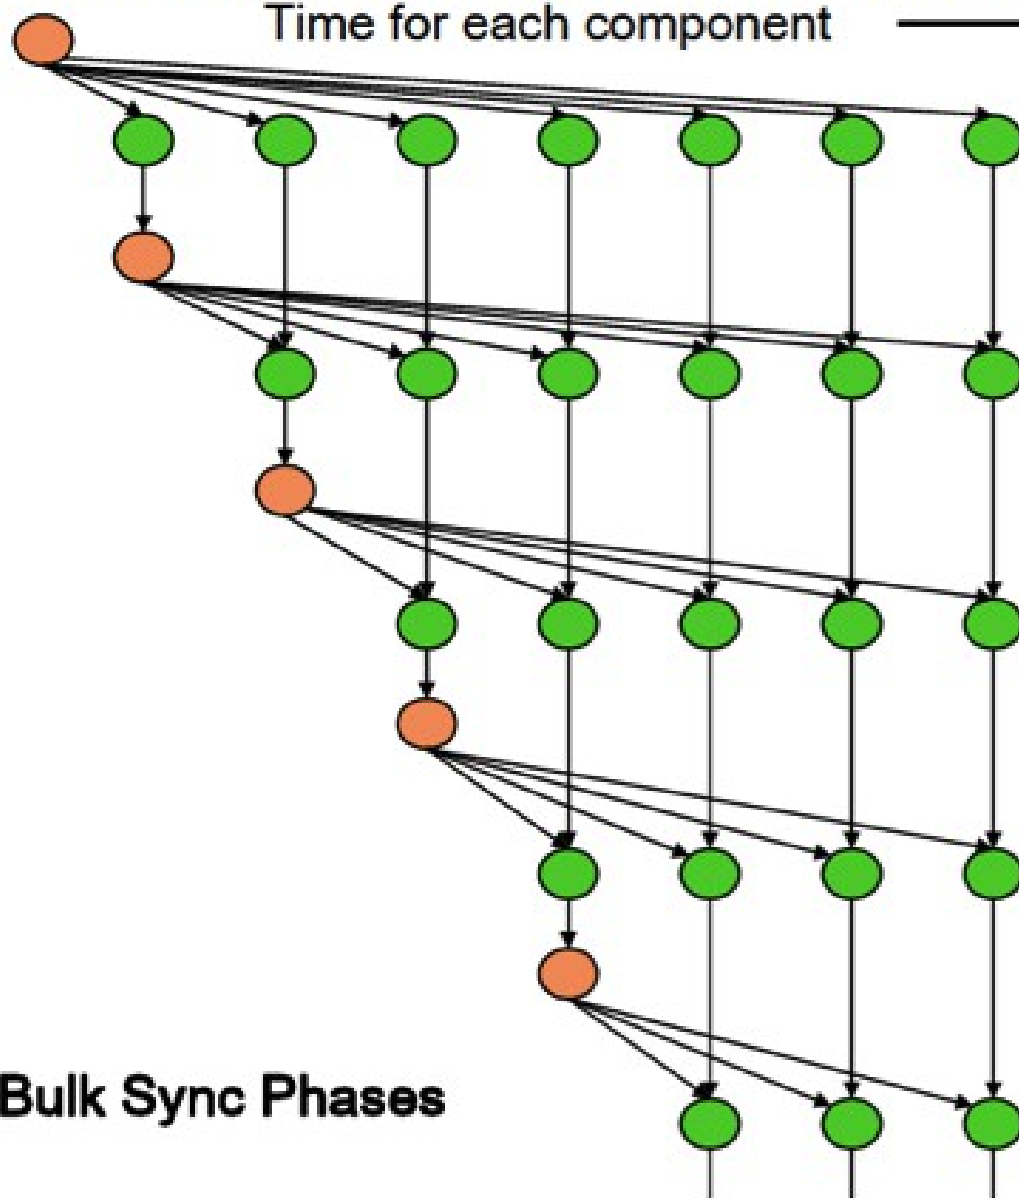
Those new algorithms need new kernels and rely on efficient scheduling algorithms.

LU Timing Profile (4 processor system)

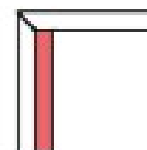
Threads – no lookahead



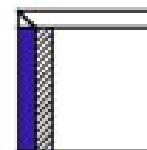
Time for each component →



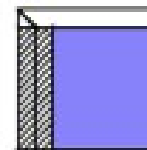
DGETF2



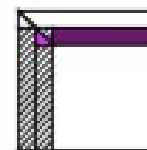
DLSWP



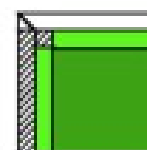
DLSWP



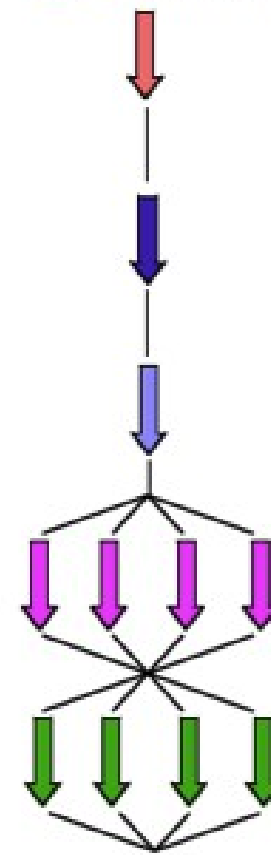
DTRSM



DGEMM

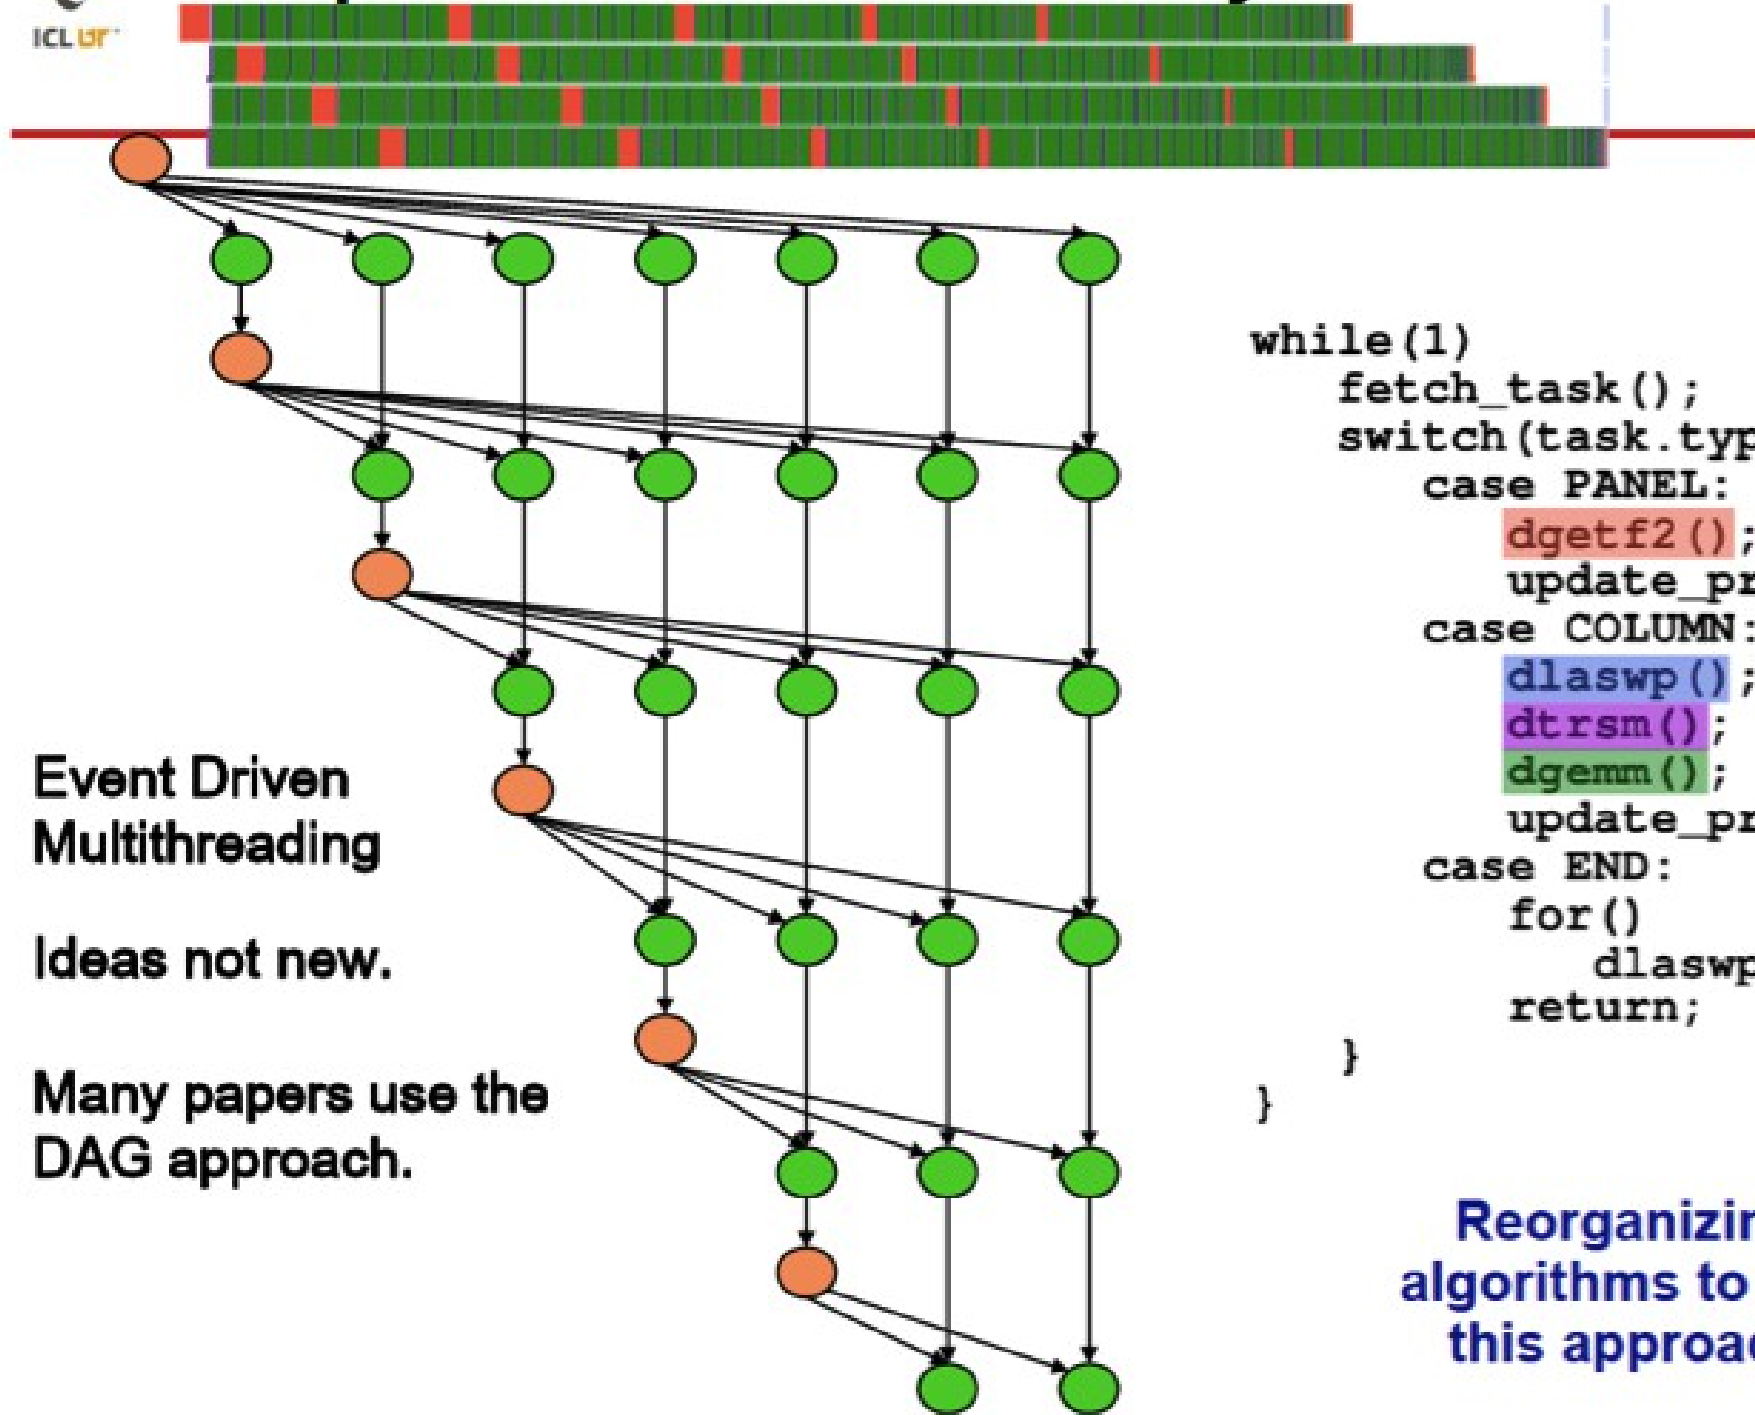


- DGETF2
- DLASWP(L)
- DLASWP(R)
- DTRSM
- DGEMM



Bulk Sync Phases

Adaptive Lookahead - Dynamic



Event Driven
Multithreading

Ideas not new.

Many papers use the
DAG approach.

```
while(1)
  fetch_task();
  switch(task.type) {
    case PANEL:
      dgetf2();
      update_progress();
    case COLUMN:
      dlaswp();
      dtrsm();
      dgemm();
      update_progress();
    case END:
      for()
        dlaswp();
      return;
  }
}
```

Reorganizing
algorithms to use
this approach

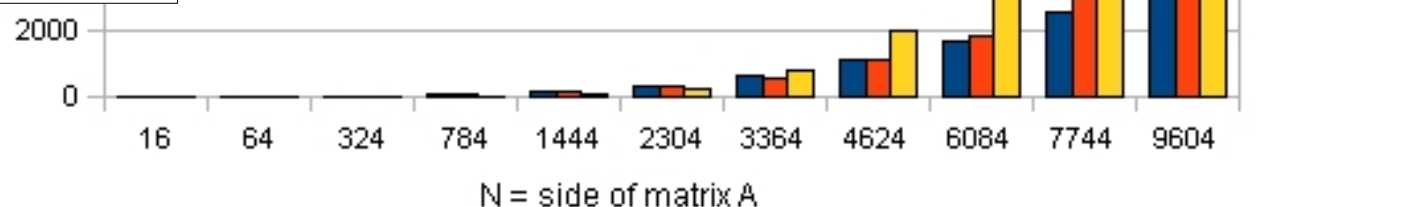
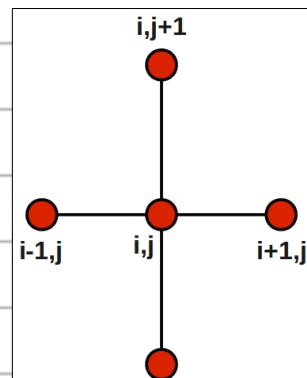
Sparse Linear Algebra

```

-4 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0
1 -4 1 0 1 0 0 0 0 0 0 0 0 0 0 0
0 1 -4 1 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 -4 1 0 1 0 0 0 0 0 0 0 0 0
0 1 0 1 -4 1 0 1 0 0 0 0 0 0 0 0
0 0 1 0 1 -4 1 0 1 0 0 0 0 0 0 0
0 0 0 1 0 1 -4 1 0 1 0 0 0 0 0 0
0 0 0 0 1 0 1 -4 1 0 1 0 0 0 0 0
0 0 0 0 0 1 0 1 -4 1 0 1 0 0 0 0
0 0 0 0 0 0 1 0 1 -4 1 0 1 0 0 0
0 0 0 0 0 0 0 1 0 1 -4 1 0 1 0 0
0 0 0 0 0 0 0 0 1 0 1 -4 1 0 1 0
0 0 0 0 0 0 0 0 0 1 0 1 -4 1 0 1
0 0 0 0 0 0 0 0 0 0 1 0 1 -4 1 0
0 0 0 0 0 0 0 0 0 0 0 1 0 1 -4

```

Iterative vs. Direct Solve

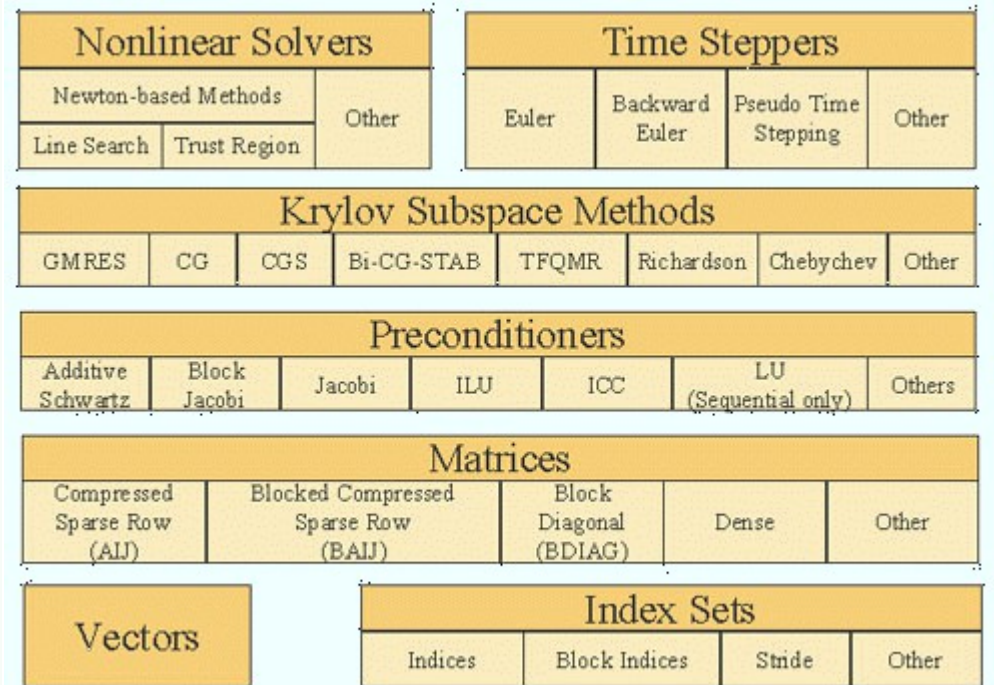
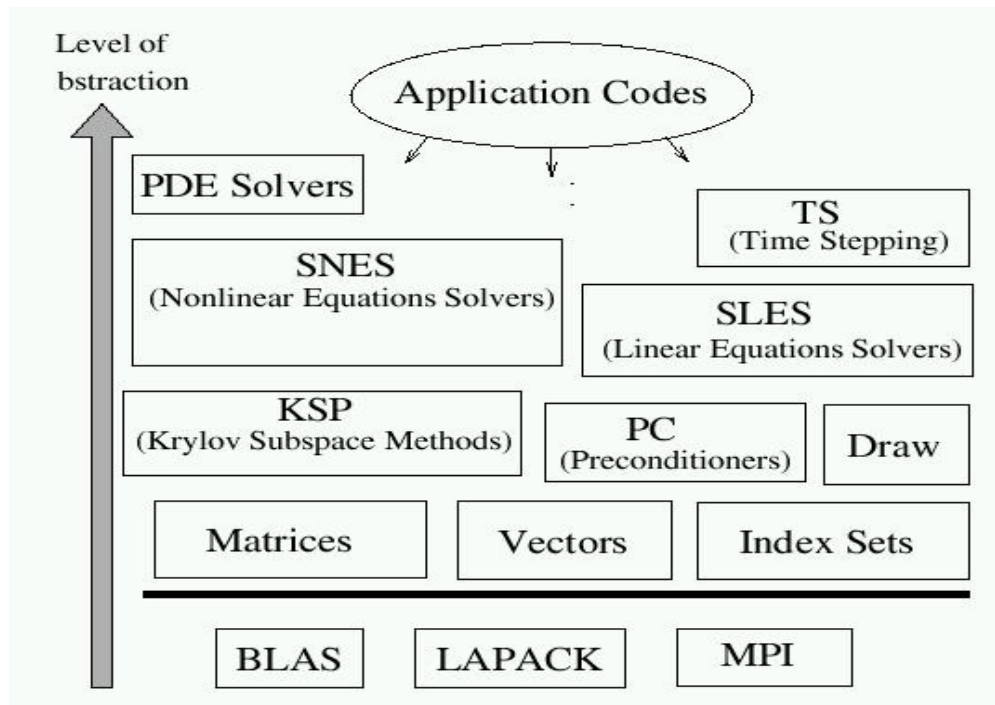


- PETSc: <http://www.mcs.anl.gov/petsc>

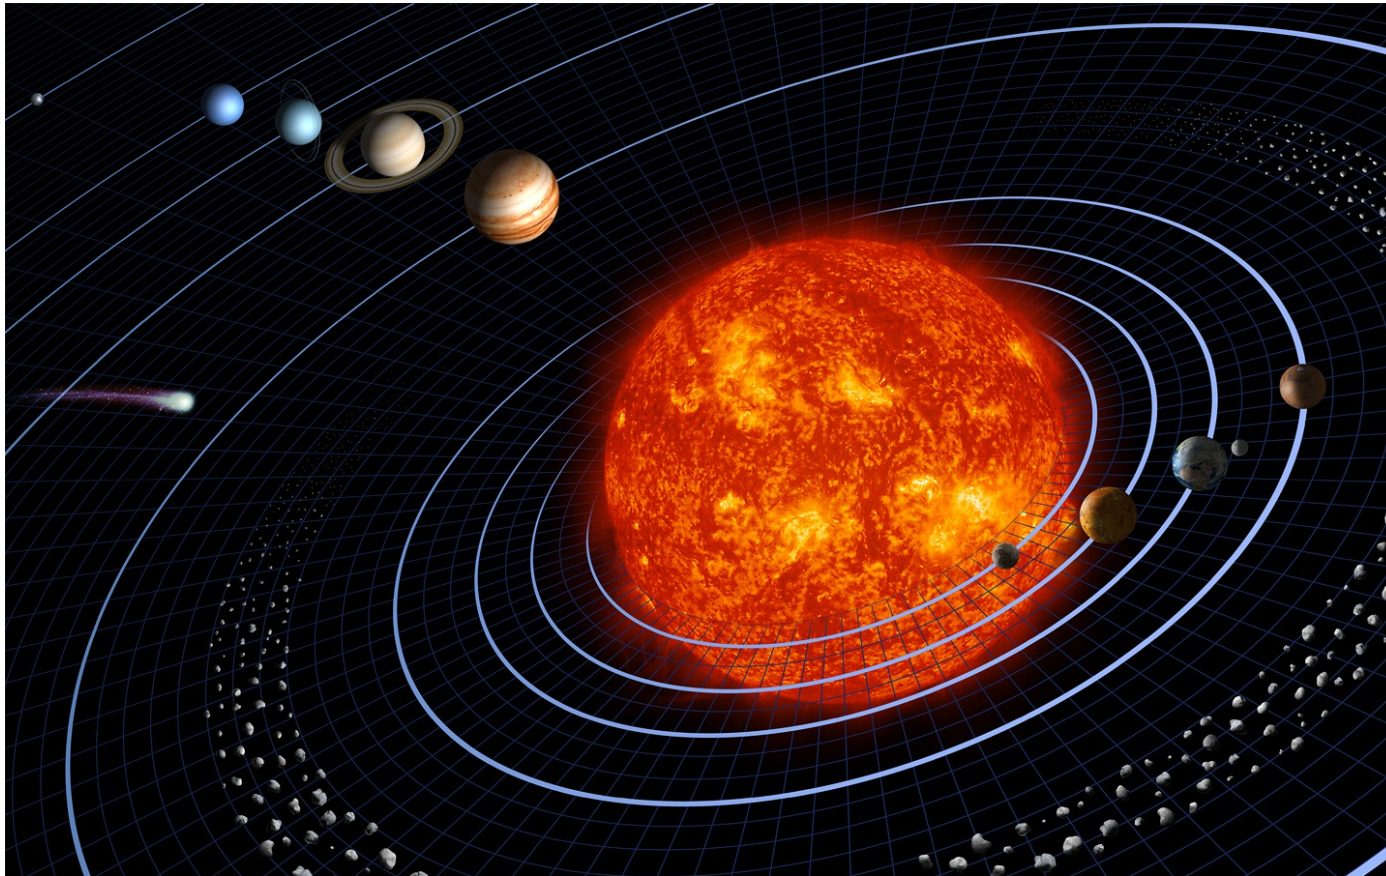
Sparse Matrix Solver: PETSc



- <http://www.mcs.anl.gov/petsc/>
- Threaded – multi-core
- MPI – distributed memory
- GPU – many-core



N-Body Problems



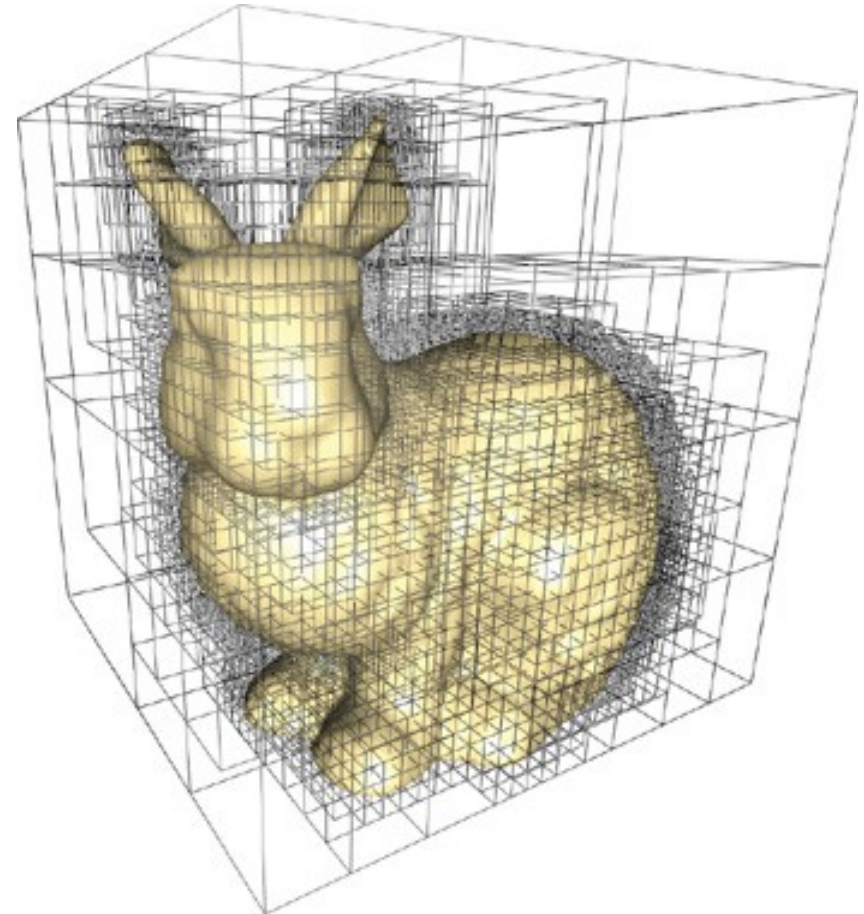
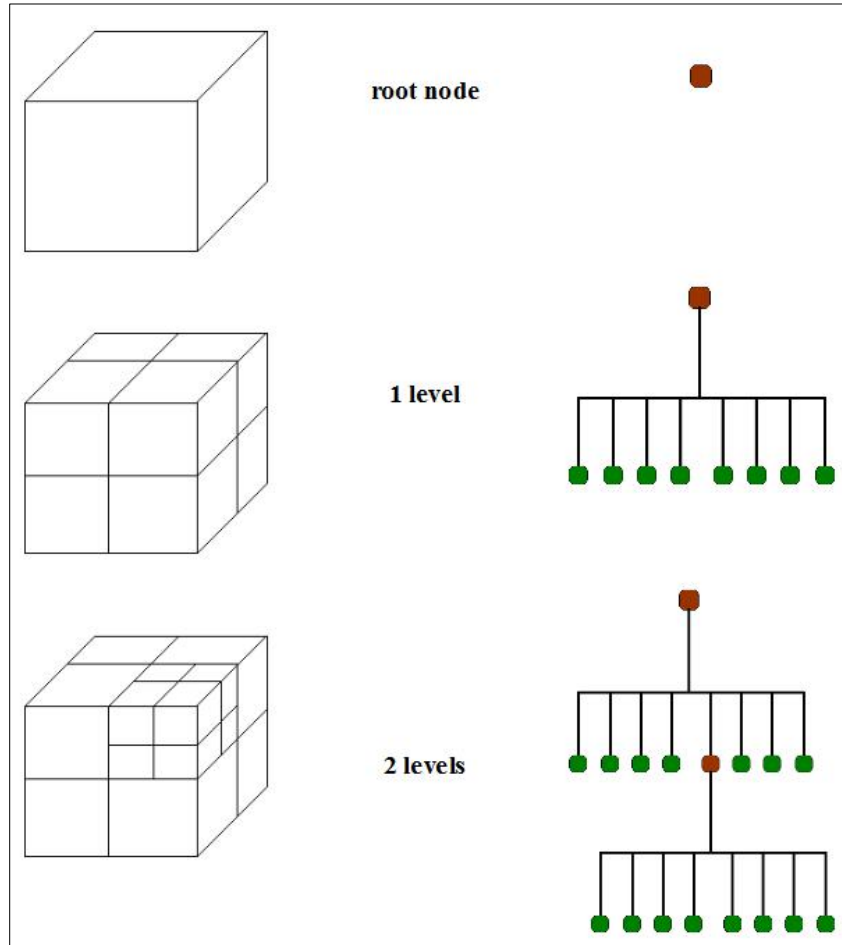
An N-body simulation numerically approximates the evolution of a system of bodies in which each body continuously interacts with every other body

OpenMM N-body library

- Aimed at molecular mechanics (MM)
- Runs on CPUs, GPUs etc
- Easy to use from Python, C/C++, ...

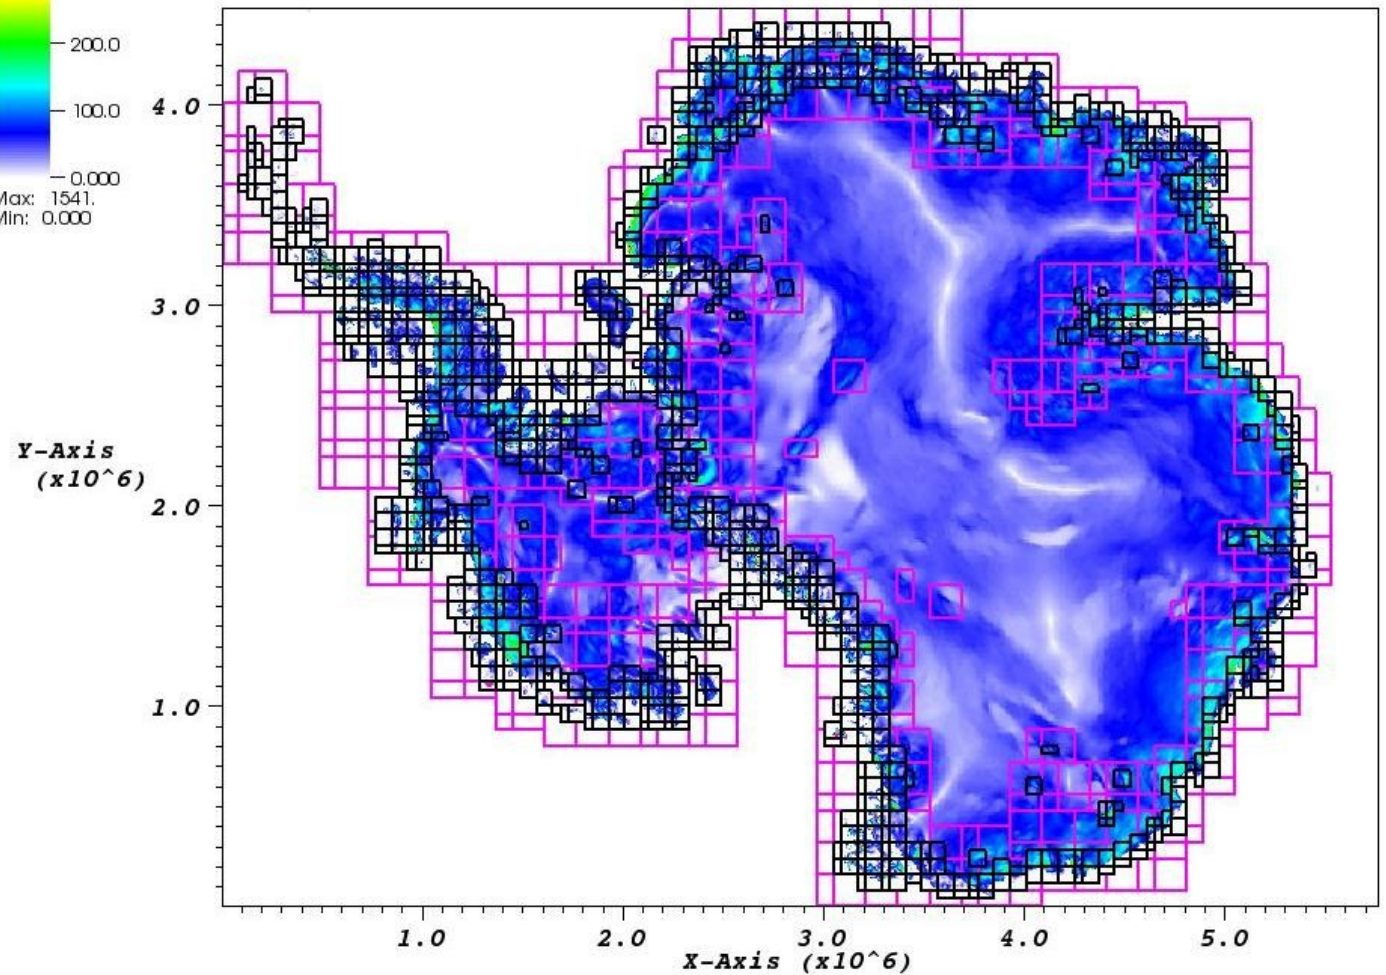
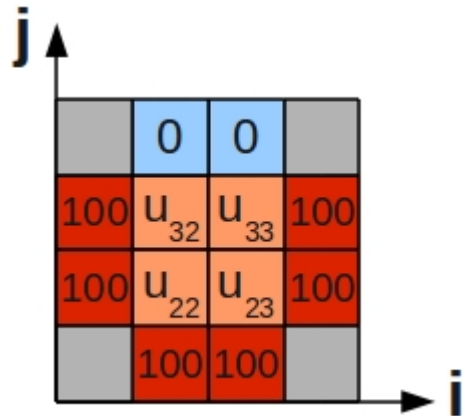
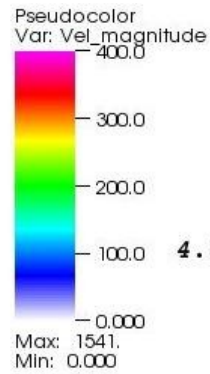
```
pdb = PDBFile('input.pdb')  
forcefield = ForceField('amber99sb.xml', 'tip3p.xml')  
system = forcefield.createSystem(pdb.topology, nonbondedMethod=PME,  
nonbondedCutoff=1*nanometer, constraints=HBonds)  
integrator = LangevinIntegrator(300*kelvin, 1/picosecond, 0.002*picoseconds) simulation =  
Simulation(pdb.topology, system, integrator) simulation.context.setPositions(pdb.positions)  
simulation.minimizeEnergy() simulation.reporters.append(PDBReporter('output.pdb', 1000))  
simulation.step(10000)
```


Tree methods: Fast Multipole Method (FMM)



FMM in PETSc: http://barbagroup.bu.edu/Barba_group/PetFMM.html

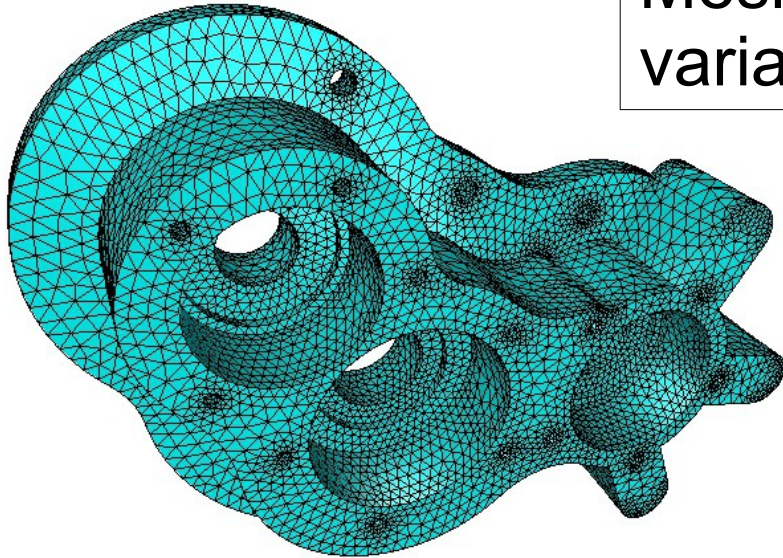
Structured Grids



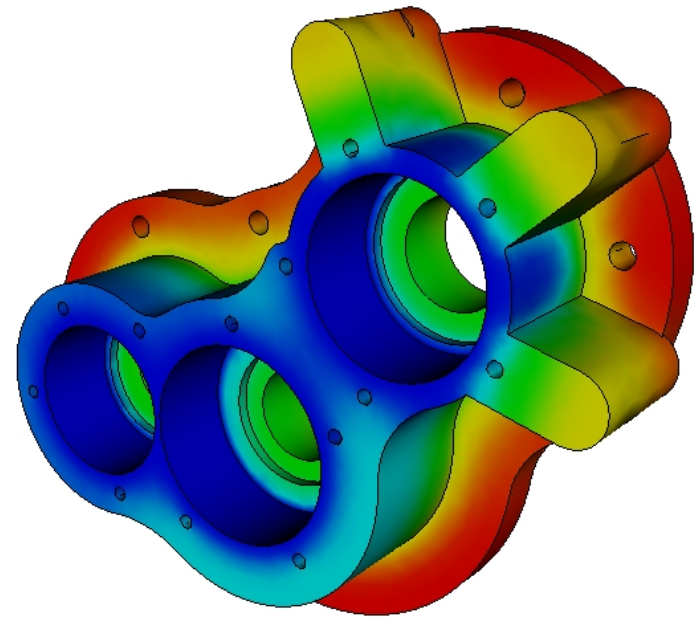
search: CHOMBO, OPS

Unstructured Grids

Mesh assembled from
variable sized polygons



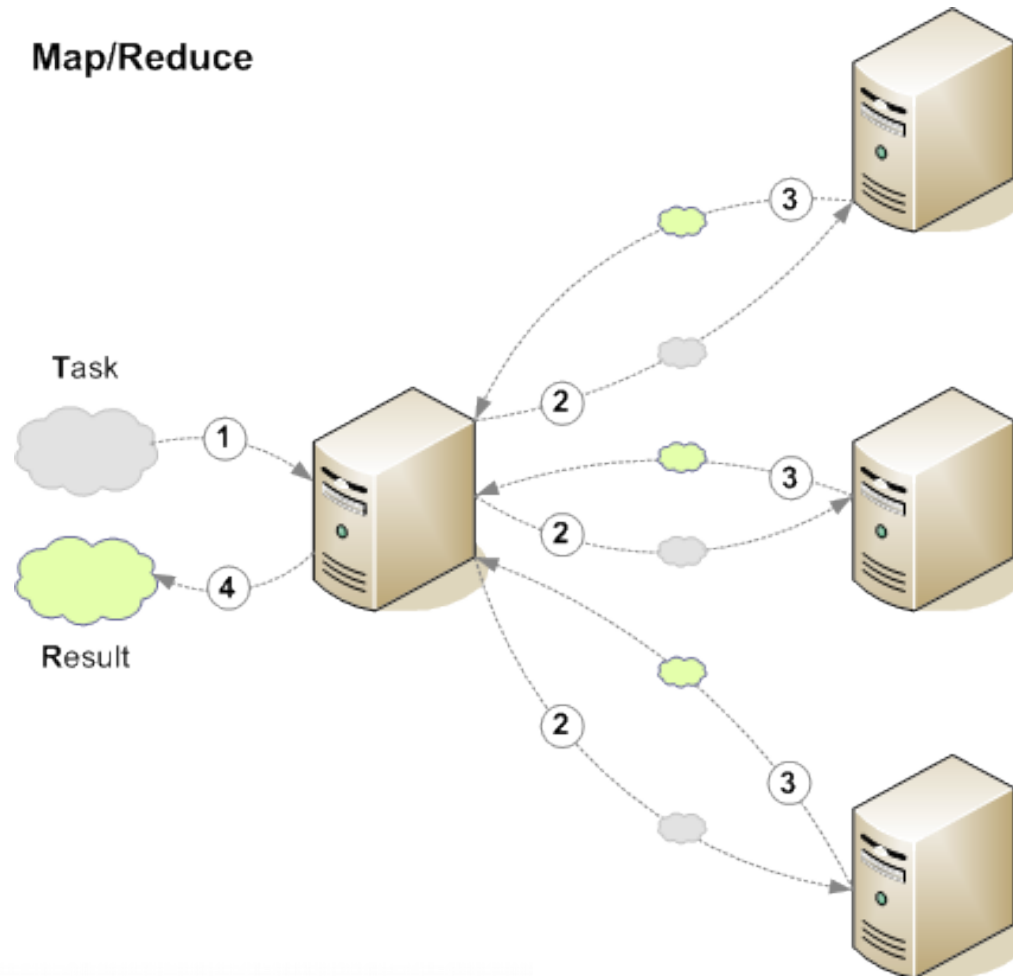
Computational expense
can be focused on areas
of interest



ELMER: <http://www.csc.fi/english/pages/elmer>
OP2: <http://www.oerc.ox.ac.uk/projects/op2>

MapReduce

Map/Reduce



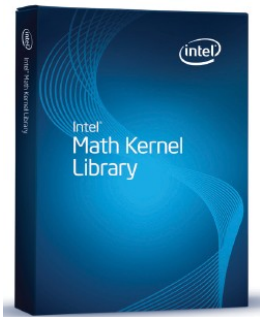
<http://hadoop.apache.org/>

Others

QuantLib

nag[®] *Results Matter. Trust NAG.*

GSL: <http://www.gnu.org/software/gsl/>



NVIDIA[®]

<http://developer.nvidia.com/>

© Simon McIntosh-Smith,
Gethin Williams



Recap: Libraries..

- Are typically well implemented & tested.
- Offer us parallel implementations of common computational kernels.
- Can fast-track us to real-world applications.
- Provide an abstraction between your application and the details of the implementation.
- Hard (impossible?) to beat for performance!

Don't reinvent the wheel!

